

In Practice

- <https://www.microsoft.com/en-us/springfield/>
- <http://www.economist.com/news/science-and-technology/21647269-automating-search-loopholes-software-hacking-hackers>
- <https://www.wired.com/2016/08/security-bots-show-hacking-isnt-just-humans/>
- <http://www.digitaltrends.com/computing/darpa-mayhem-bot/>
- ...

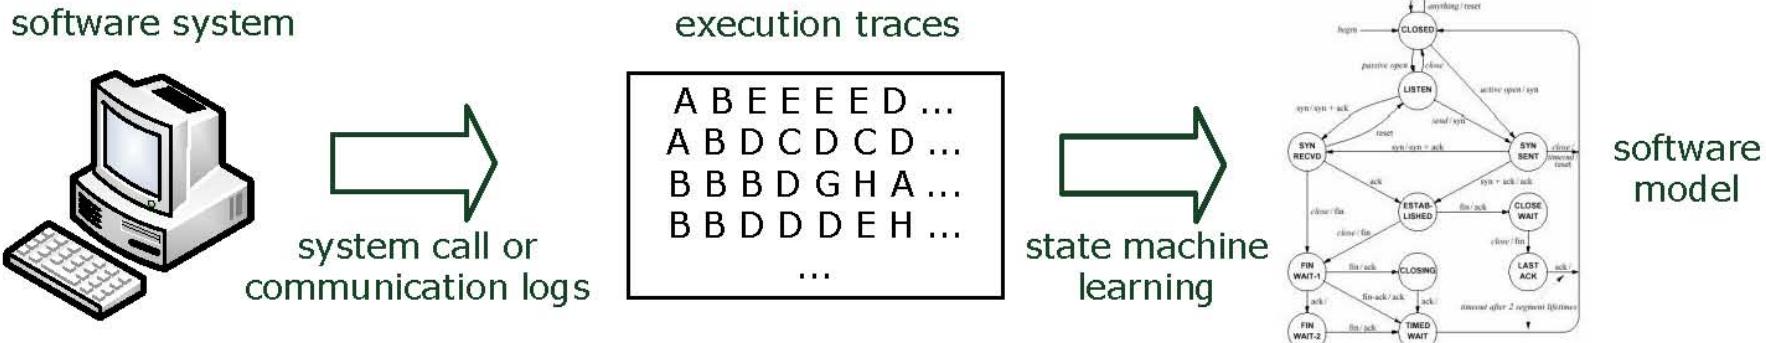
Teminology

- Black-box (fuzz/random) testing:
 - Not interpreting the software, using it as a “black-box”
 - Boundary value analysis
 - Brute-force
- White-box (fuzz) testing:
 - Symbolic Execution
 - Unit Testing
- Grey-box testing:
 - Some parts are black-box, other white-box
 - Concolic, DART (?)
 - Regression testing

Today: State Machine Learning

Myhill-Nerode
language Hankel matrix
graph coloring
active and passive learning

System identification and analysis



- Software leaves **traces**
- A state machine is a **logical model** describing these traces
 - Classification – is a new trace generated by the same software?
 - Prediction – what trace is most likely to occur next?
 - **Analysis** – is the software deadlock-free, secure, **malicious**?

System identification and analysis

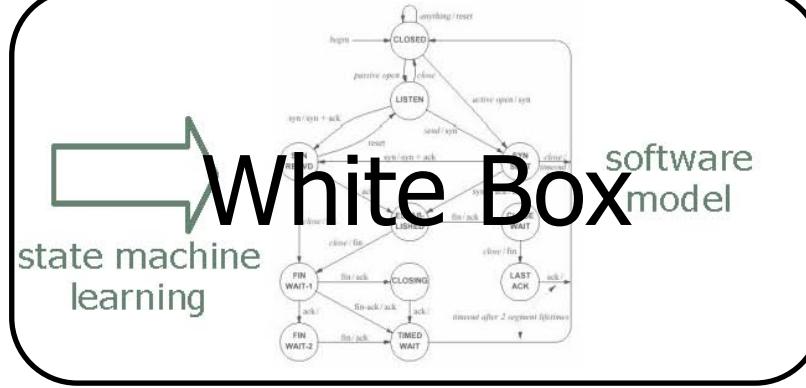
Black Box

execution traces

```
A B E E E D ...  
A B D C D C D ...  
B B B D G H A ...  
B B D D D E H ...  
...
```

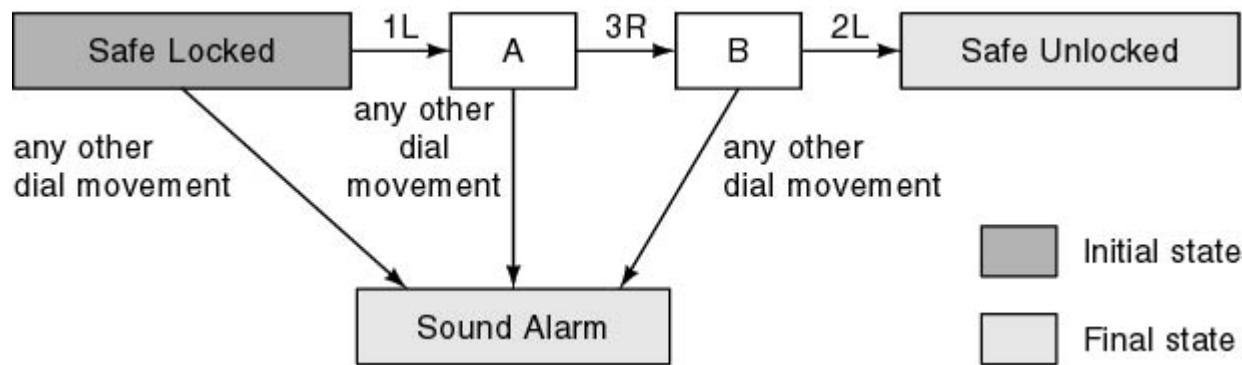
state machine
learning

White Box
software
model

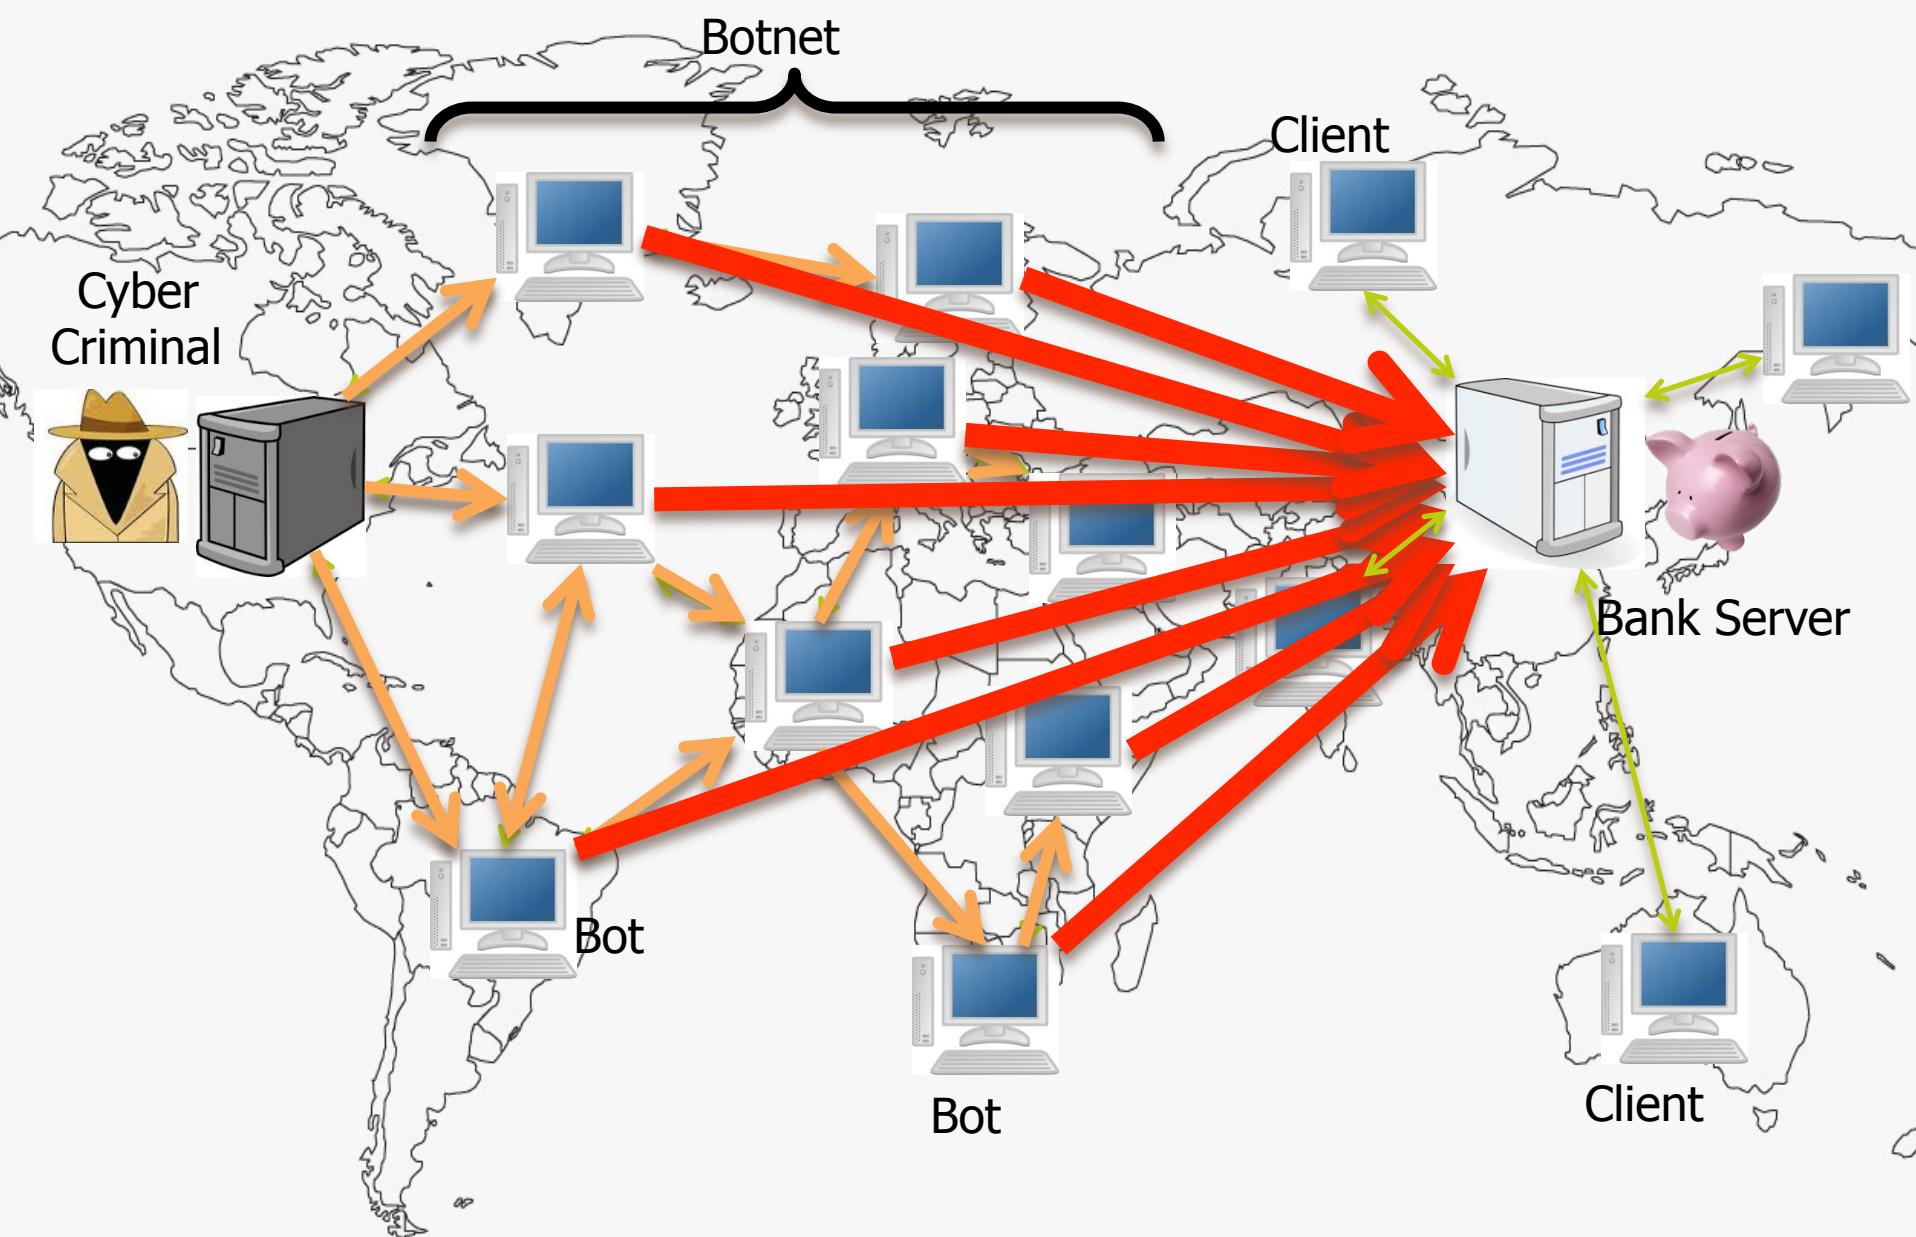


State machines

- State machines are key models for the design and analysis of software systems
- They model the states of a system and the transitions between them:

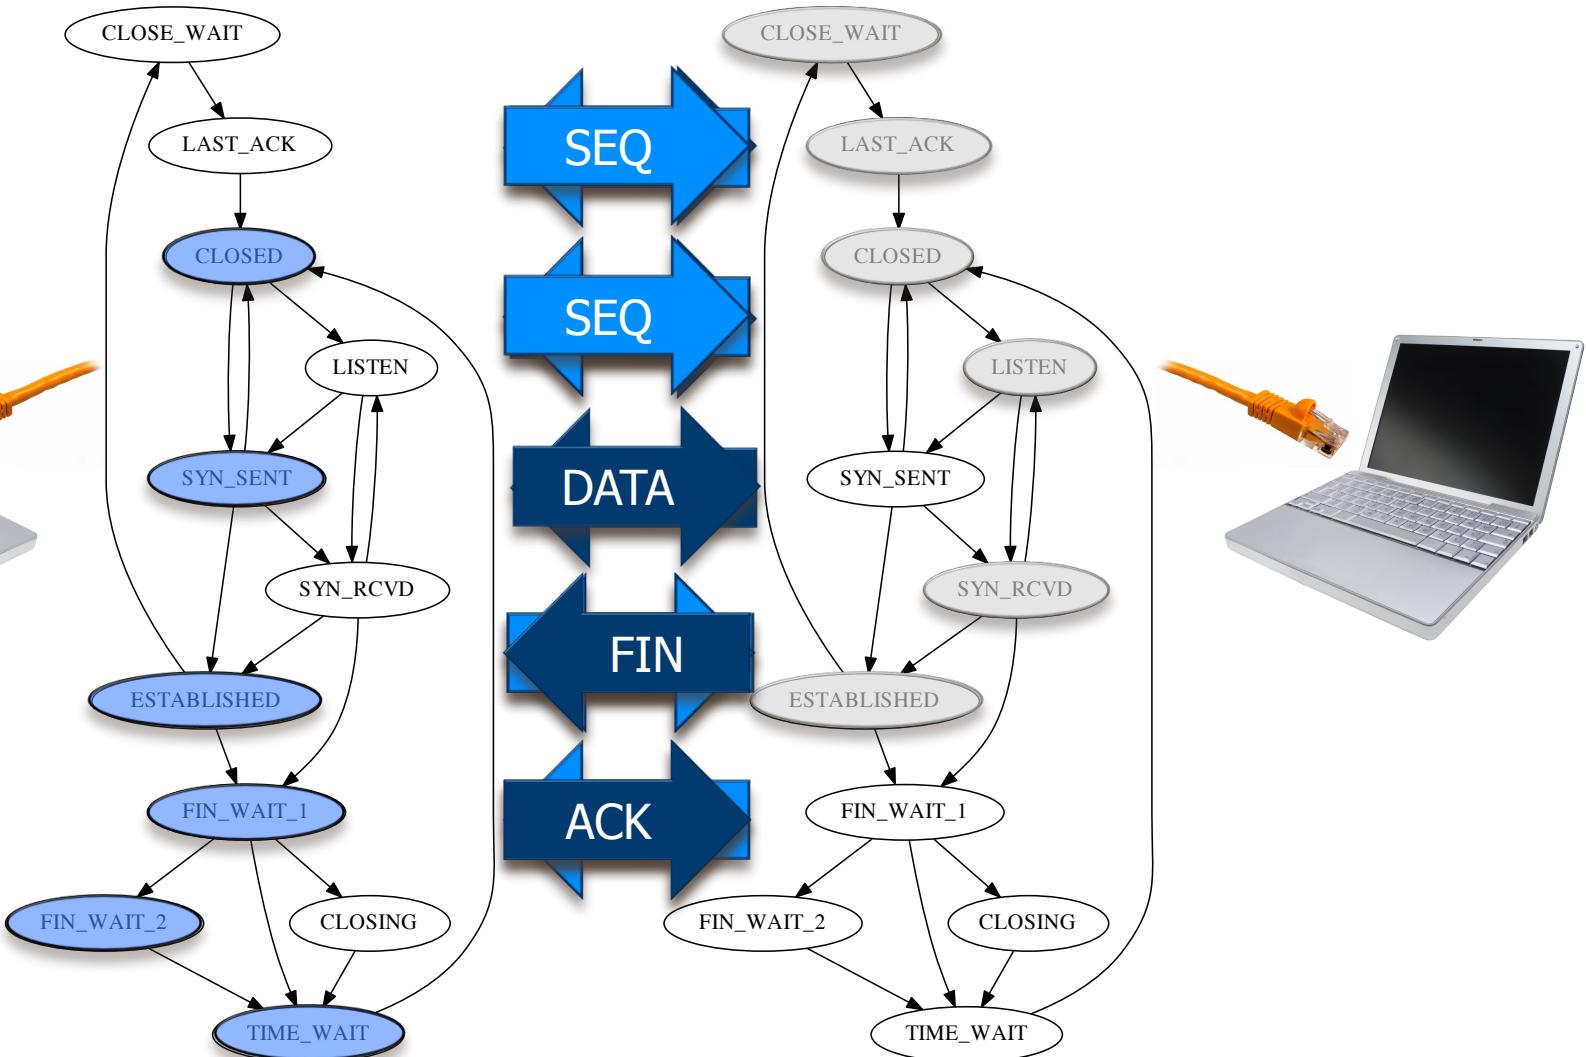


- Used for visualization, verification, testing, control, ...



We observe internet traffic

- SYN, SYN-ACK, ACK, SEQ, ACK, SEQ, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
-



Learning (reverse-engineering)

- Execution logs form a **big data** set from which can be used to gain information about a system or protocol

This can help to

- analyze **your own code** and hunt for bugs, or
- reverse-engineer **someone else's unknown protocol**, eg. a botnet, to fingerprint or to analyze (and attack)

We observe internet traffic

- SYN, SYN-ACK, ACK, SEQ, ACK, SEQ, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
-

First: make abstract

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

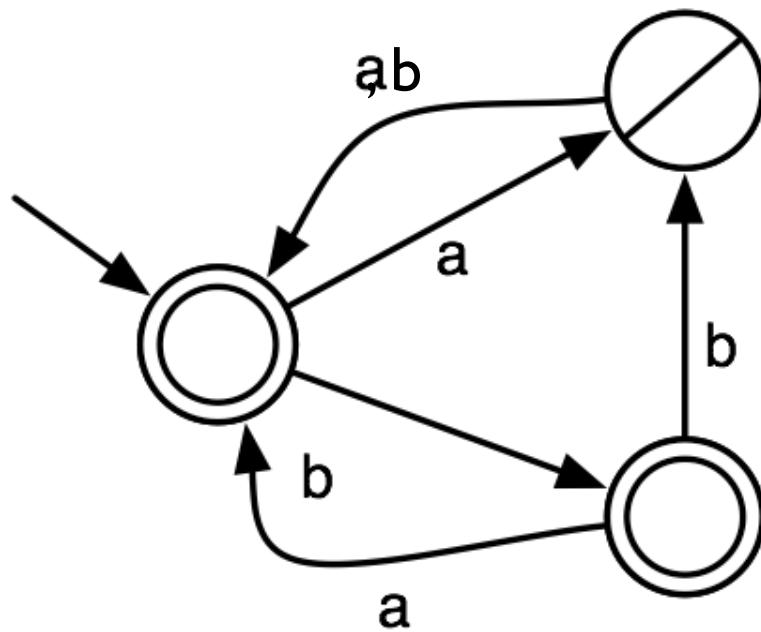
Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Deterministic finite state automata (DFA)



$(a,-), (ab,+), (aa,+), (b,+), (bb,-), (bab,+)$

State machine learning

- Input:
 - Data, e.g., (a,-),(ab,+),(aa,+),(b,+),(bb,-),(bab,+)
- Goal: Find a smallest DFA that is **consistent** with the data
- Size is measured by the number of states/transitions
 - Occam's razor
- Consistency is:
 - Accepts all positive, rejects all negative

Computational complexity

- NP
 - Class of problems **verifiable** in **polynomial time**
- NP-hard
 - **At least as difficult to solve as any other problem in NP**
 - Solving in polynomial time is widely believed to be impossible (even with quantum computers)
- NP-complete
 - NP-hard and in NP, e.g., satisfiability, traveling salesman
- Undecidable
 - Impossible to solve using a computer, equivalent to the halting problem

Example

accept:

1 0 0 0 0 0 0 0
0 1 1 0 0 0 1 0 1 0 1 0 1 0 0
0 0
1 0 1 0 0
1 0 0
1 0 1
1 0 1
1 0 0 0 0 0 0

reject:

0 1 0 0 0
1 1 0 0 0 0 0 0 1 0 0
0 0 0 0 0 1
1 0 0 0 0 1 0 1 0 0 0 0 1 1 1 0 0 0
0 1 1 1 0 1 0 0 0 0 0 0

What is the smallest state machine that accepts the top and rejects the bottom strings?

Guess: how many states are needed?

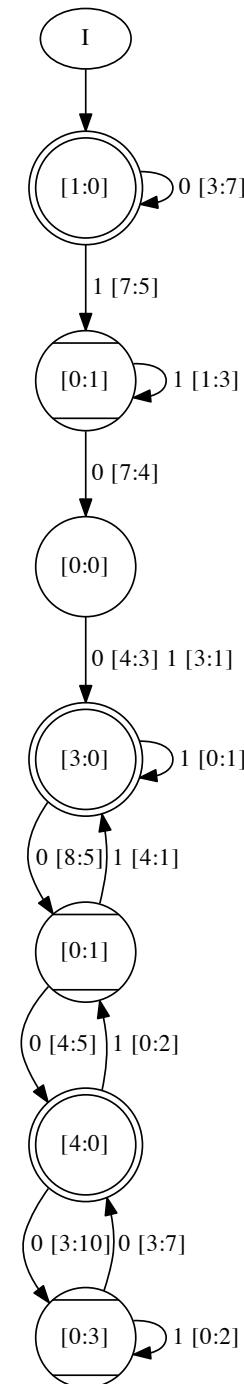
Example

accept:

```
1 0 0 0 0 0 0 0  
0 1 1 0 0 0 1 0 1 0 1 0 1 0 0  
0 0  
1 0 1 0 0  
1 0 0  
1 0 1  
1 0 1  
1 0 0 0 0 0 0
```

reject:

```
0 1 0 0 0  
1 1 0 0 0 0 0 0 1 0 0  
0 0 0 0 0 1  
1 0 0 0 0 1 0 1 0 0 0 0 1 1 1 0 0 0  
0 1 1 1 0 1 0 0 0 0 0 0 0
```



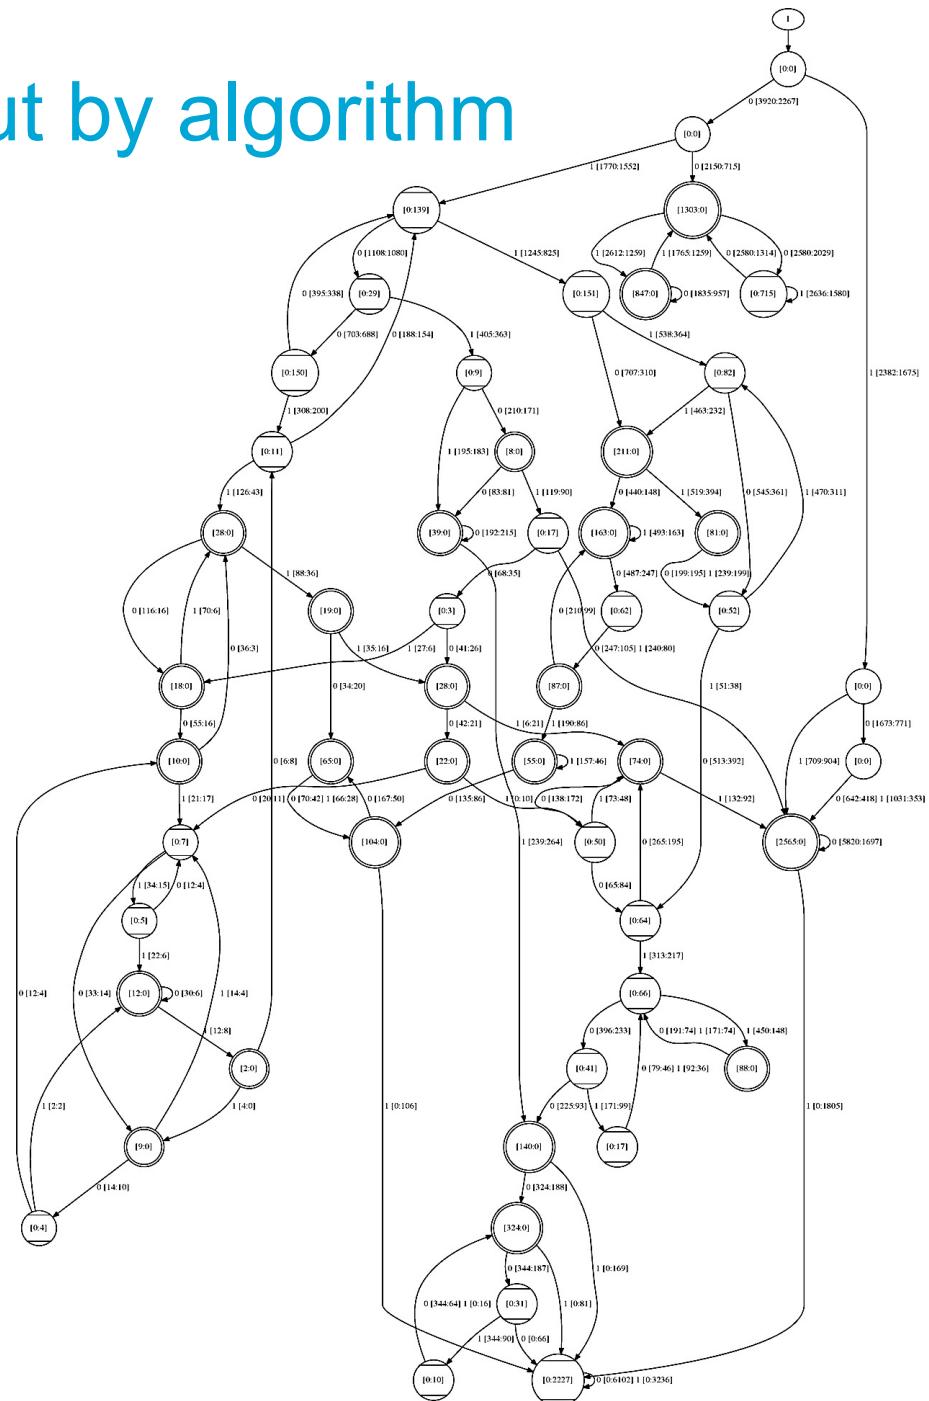
Example output by algorithm

accept:

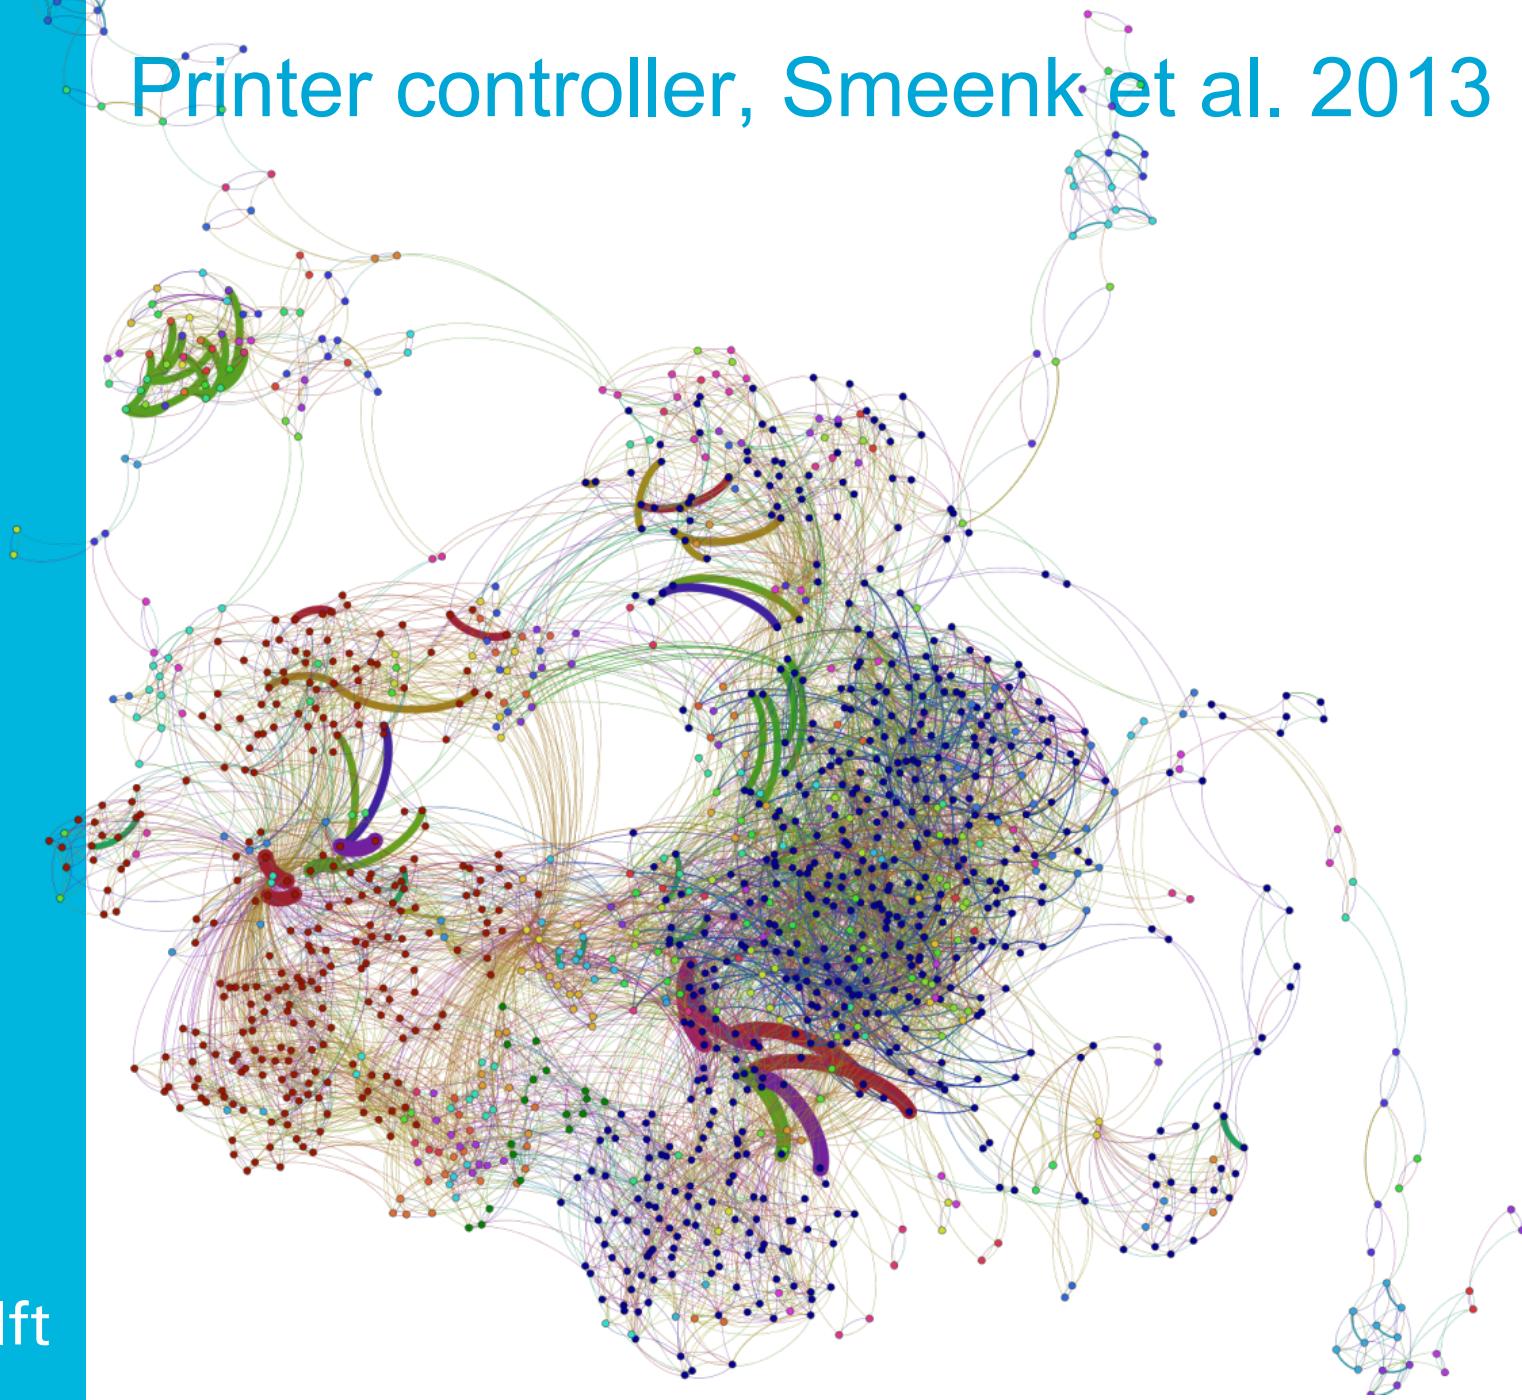
```
1 0 0 0 0 0 0 0 0
0 1 1 0 0 0 1 0 1 0 1 0 1 0 1 0 0
0 0
1 0 1 0 0
1 0 0
1 0 1
1 0 1
1 0 1
1 0 0 0 0 0 0
```

reject:

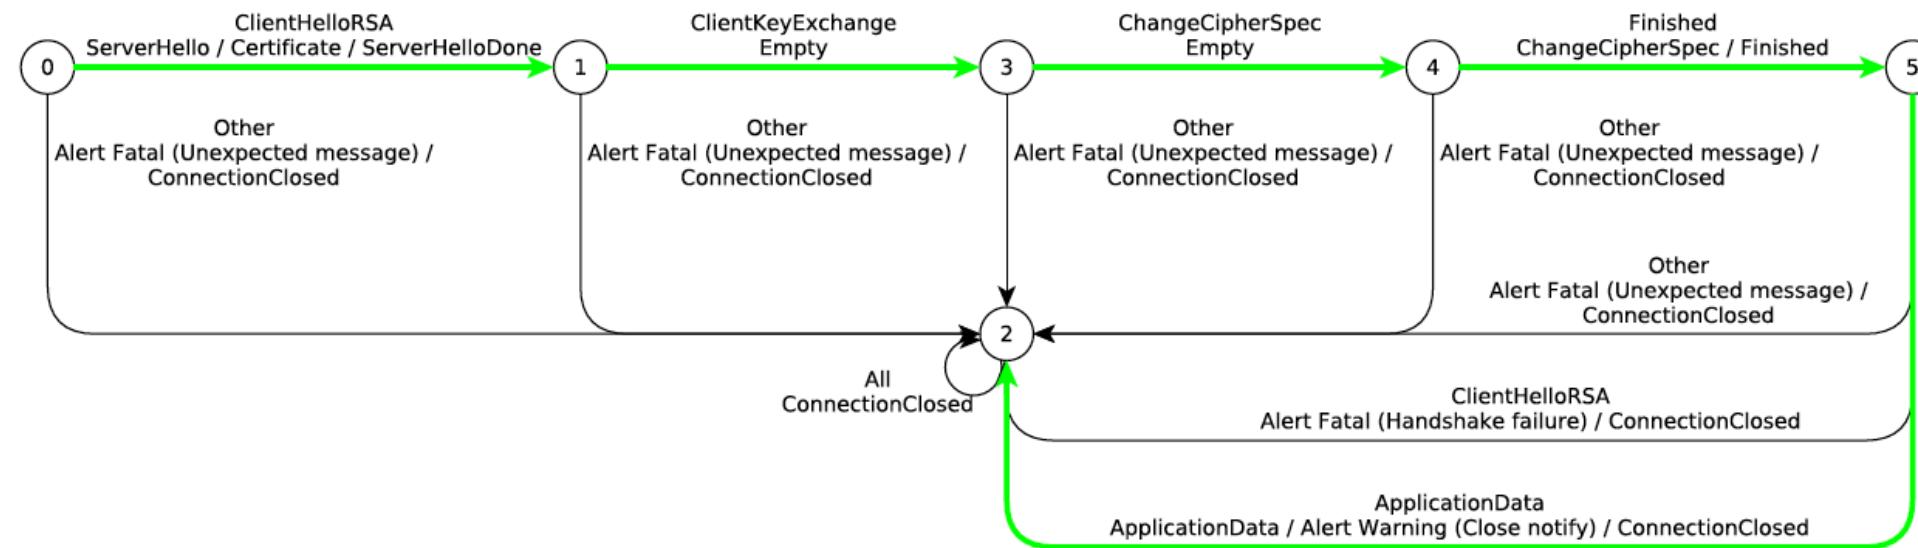
```
0 1 0 0 0
1 1 0 0 0 0 0 0 1 0 0
0 0 0 0 0 1
1 0 0 0 0 1 0 1 0 0 0 0 0 1 1 1 0 0 0
0 1 1 1 0 1 0 0 0 0 0 0 0
```



Printer controller, Smeenk et al. 2013



Use Case: TLS RSA BSAFE



Use Case: GNU TLS 3.3.8

ClientHelloRSA | ClientKeyExchange | EmptyCertificate | ChangeCipherSpec | ApplicationData
Alert Fatal (Unexpected message) / ConnectionClosed

ClientKeyExchange | EmptyCertificate | Finished | ApplicationData
Alert Fatal (Unexpected message) / ConnectionClosed

Other
Alert Fatal (Unexpected message) / ConnectionClosed

0 ServerHello Certificate / CertificateRequest / ServerHelloDone
ClientHelloRSA
ApplicationDataEmpty | HeartbeatRequest
Empty

1 ApplicationDataEmpty
Empty

ClientKeyExchange
Empty

ApplicationDataEmpty
Empty

ChangeCipherSpec
Empty

ApplicationDataEmpty
Empty

Finished
ChangeCipherSpec / Finished

ApplicationDataEmpty
Empty

HeartbeatRequest
Empty

Other
Alert Fatal (Unexpected message) / ConnectionClosed
HeartbeatRequest
Empty

ApplicationDataEmpty
Empty

EmptyCertificate
Empty

ClientKeyExchange
Empty

ClientHelloRSA
Empty

HeartbeatRequest
Empty

HeartbeatRequest
Empty

ApplicationDataEmpty |
HeartbeatRequest
Empty

Other
Alert Fatal (Unexpected message) / ConnectionClosed

ApplicationDataEmpty | HeartbeatRequest
Empty

Other
Alert Fatal (Unexpected message) / ConnectionClosed

ClientKeyExchange
Empty

ChangeCipherSpec
Empty

ApplicationDataEmpty |
HeartbeatRequest
Empty

Other
Alert Fatal (Unexpected message) / ConnectionClosed

Other
Alert Fatal (Unexpected message) / ConnectionClosed

Other
Alert Fatal (Internal error) / ConnectionClosed

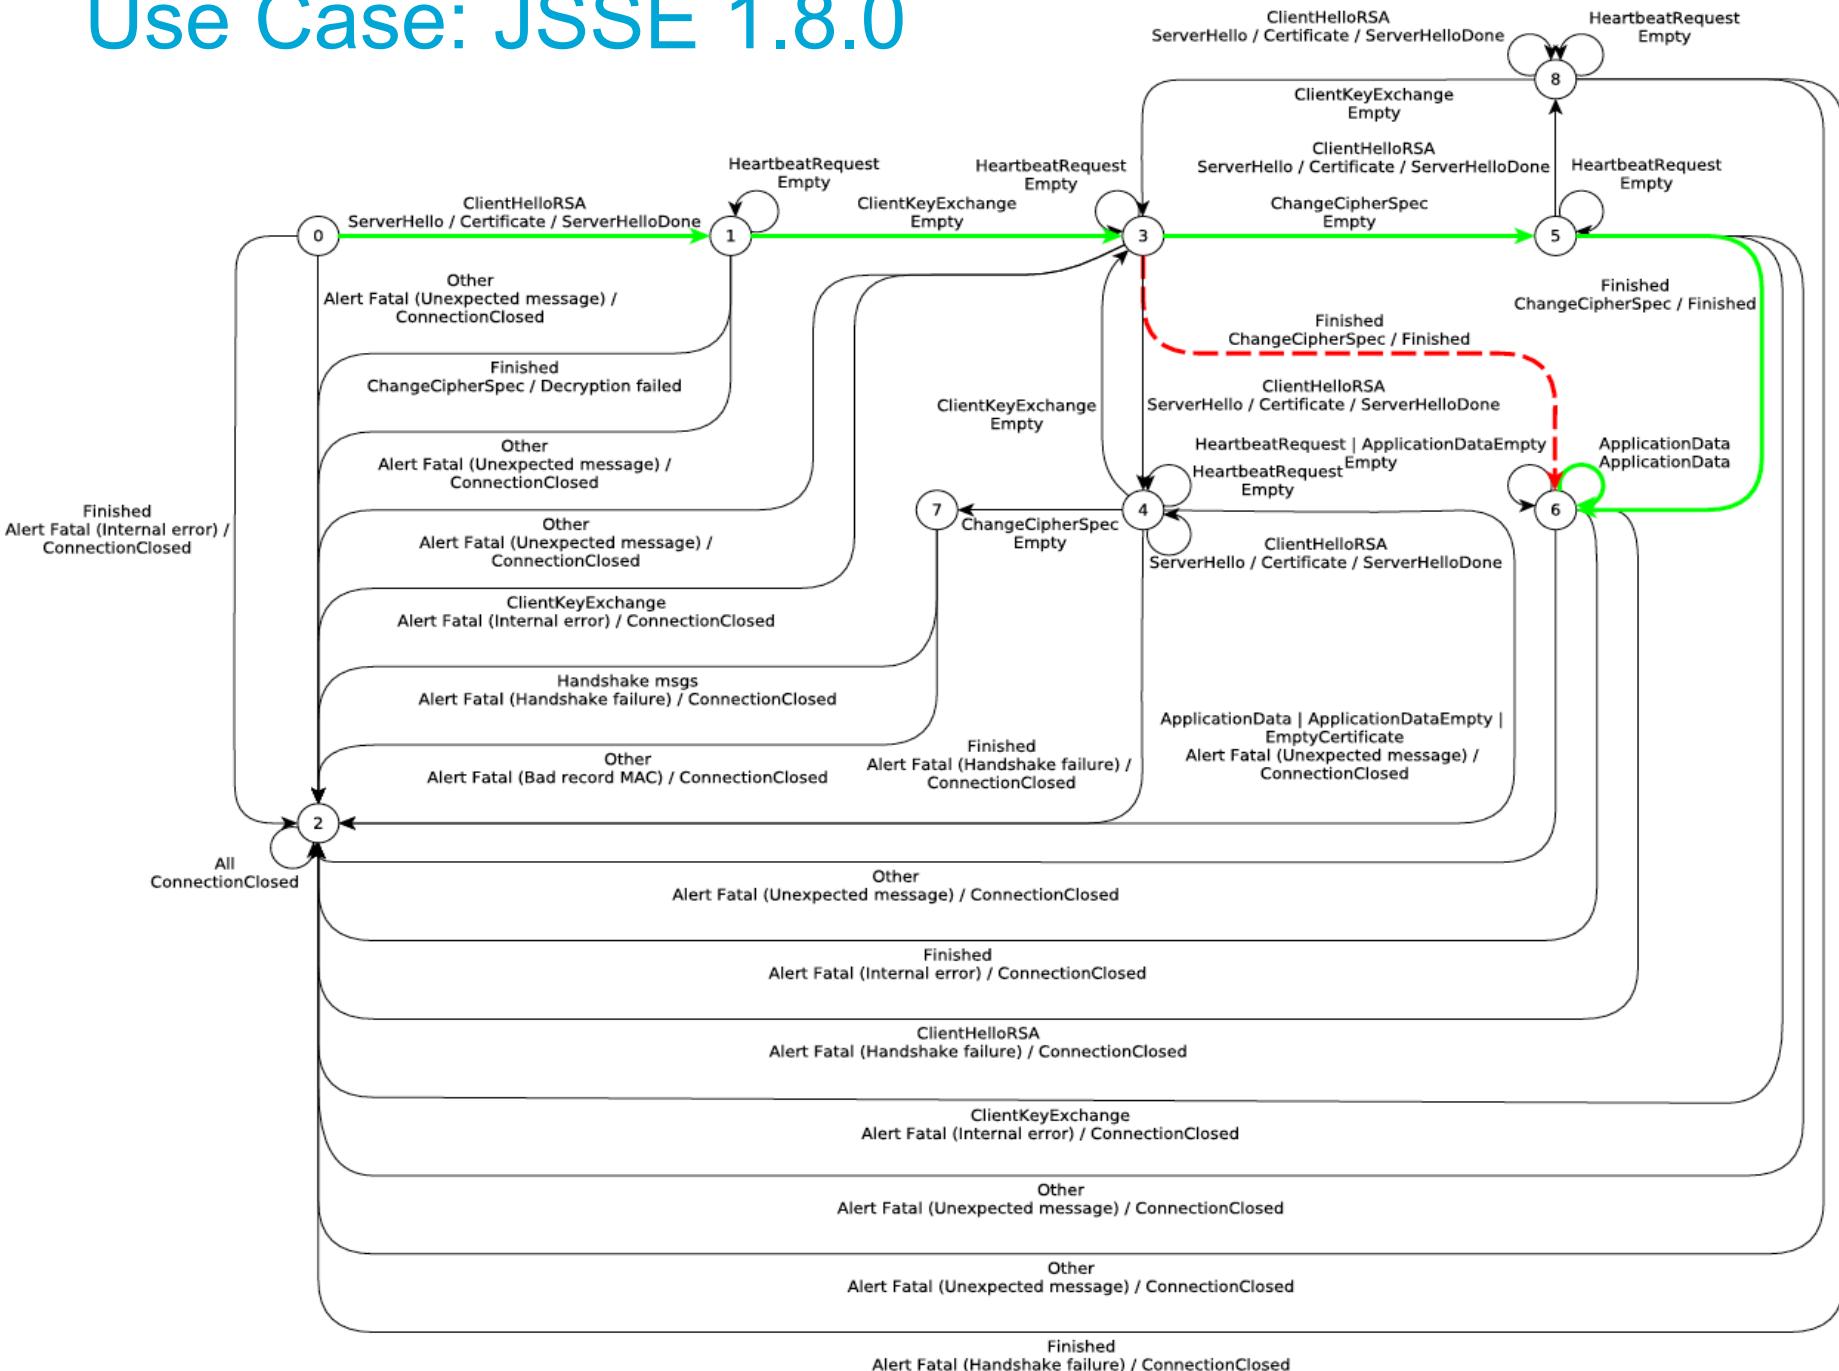
Finished
Alert Fatal (Handshake failure) / ConnectionClosed

ClientHelloRSA
Alert Fatal (Handshake failure) / ConnectionClosed

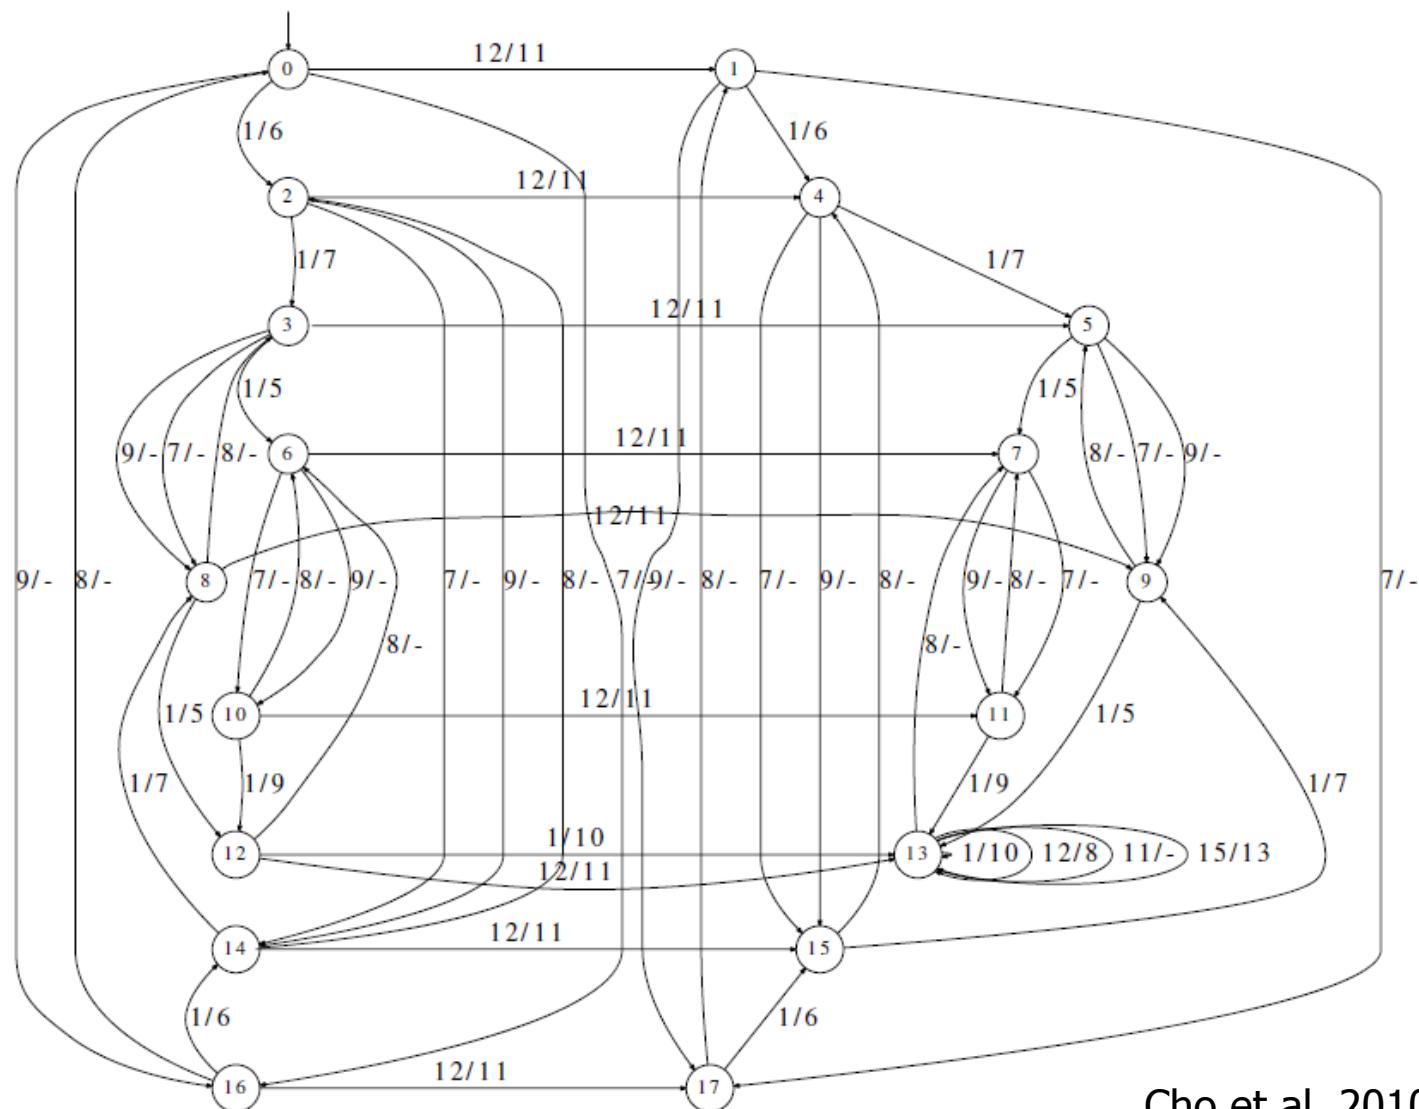
ClientKeyExchange | EmptyCertificate | ChangeCipherSpec | Finished
Alert Warning (Close notify) / ConnectionClosed

ApplicationData
ApplicationData / Alert Warning (Close notify) / ConnectionClosed

Use Case: JSSE 1.8.0

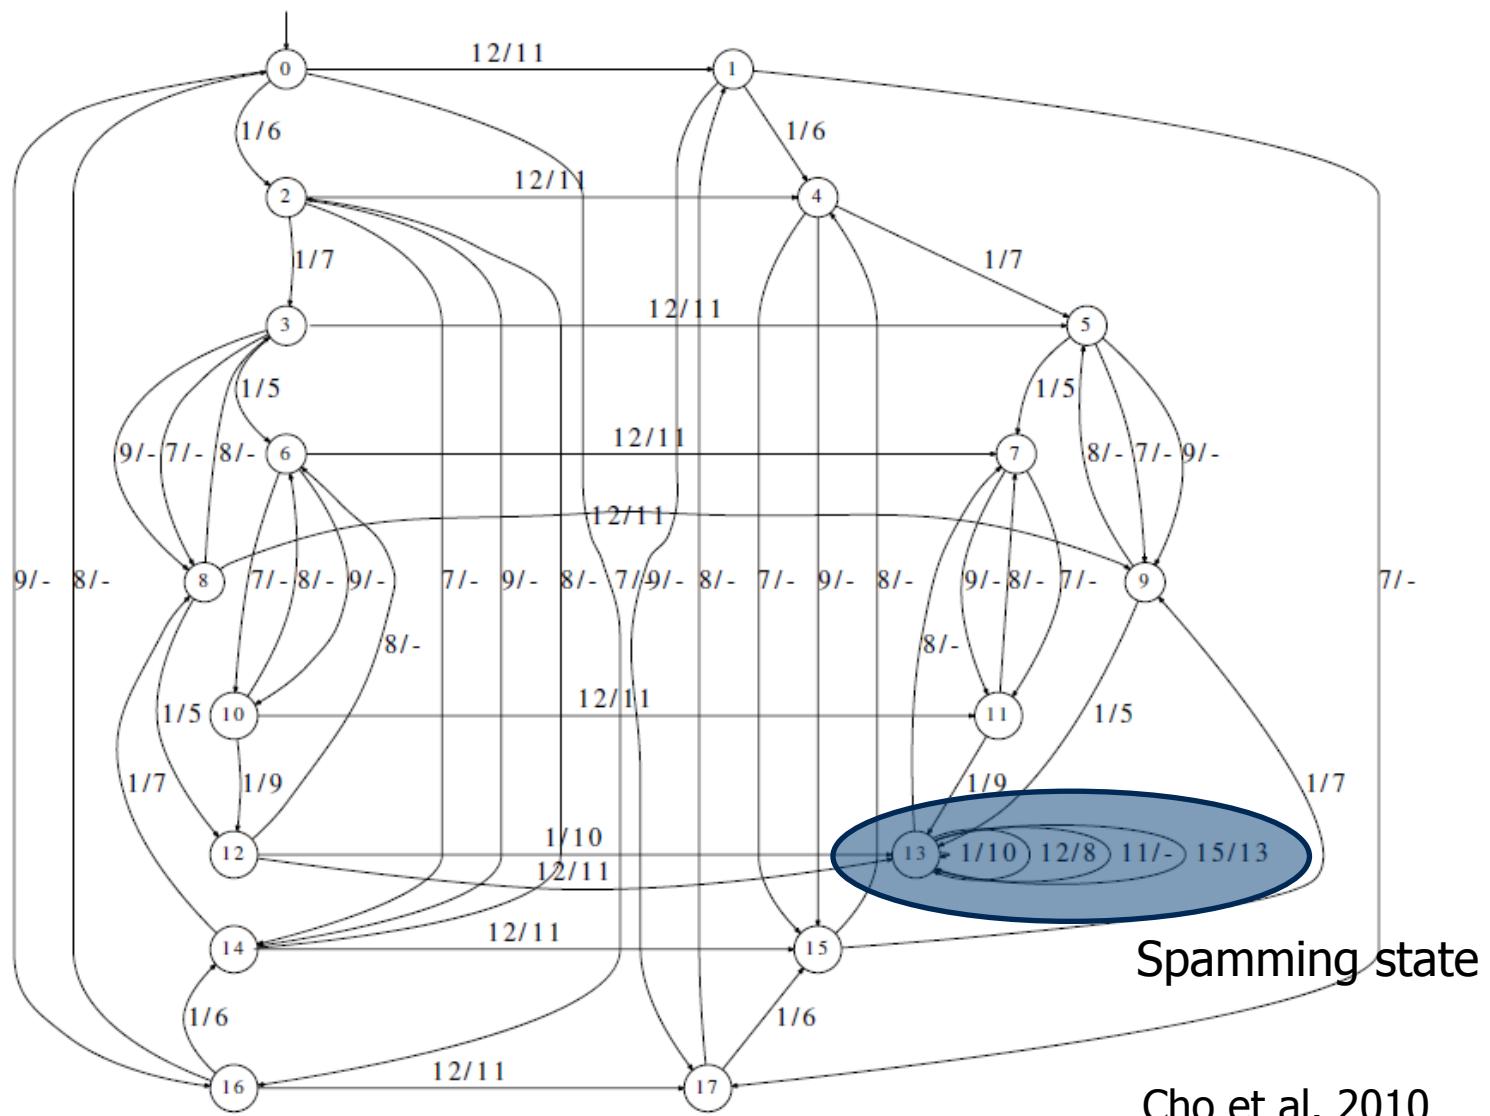


Use Case: MegaD botnet protocol



Cho et al. 2010

Use Case: MegaD botnet protocol



MYHILL-NERODE AND THE HANKEL MATRIX

Myhill-Nerode

- A **necessary** and **sufficient** condition (if and only if) for a language to be regular
 - Characterises when finite state machines exist
 - We use it for learning such state machines

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a **distinguishing extension** to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a **relation R** on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an **equivalence relation**.
- L is regular if and only if R has a **finite number of equivalence classes** (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a **distinguishing extension** to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a **relation R** on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, *the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .*

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a **distinguishing extension** to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

Visualization

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa										
ba										

- $x = aa$
- $y = ba$
- $z = \text{column}$

two strings aa and ba
and possible distinguishing extensions

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa										
ba										

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0									
ba										

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0									
ba	0									

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0	0								
ba	0	0								

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0	0	1							
ba	0	0	0							

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	Hankel matrix of language L								1 0
ab	1									1 0
ba	0									0 0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

Hankel matrix

- Rows are prefixes
- Columns are suffixes
- Row-column (x,y) is:
 - 1 if xy in L
 - 0 if xy not in L

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	x R y if row x and row y have the same color									0
ab										0
ba										0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- **It is easy to show that R is an equivalence relation.**
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a **finite number of equivalence classes** (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	1	0	1	0	1	0	1	0	1
aa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

A language is regular if and only if the Hankel matrix contains a finite number of distinct rows,
i.e., colors

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the **smallest deterministic finite state automaton** recognizing L is equal to the number of equivalence classes in R .

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	1	0	1	0	1	0	1	0	1
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	1	0	1	0	1	0	1	0
bb	0	1	0	1	0	1	0	1	0	1
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

The number of state of the smallest DFA for L is
the number of colors in the Hankel matrix

Constructing a state machine

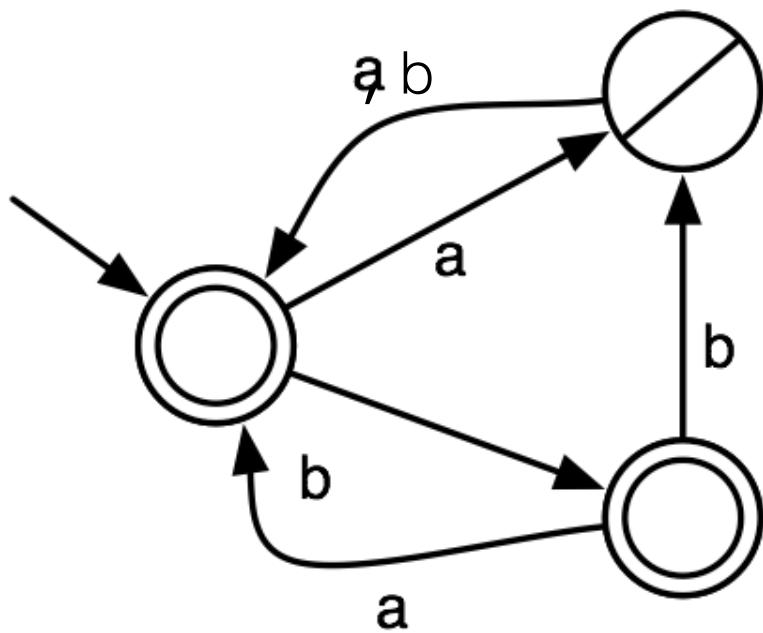
- Given a Hankel matrix for L
- What is the smallest state machine (DFA) for L ?

Q: what does this state machine look like?

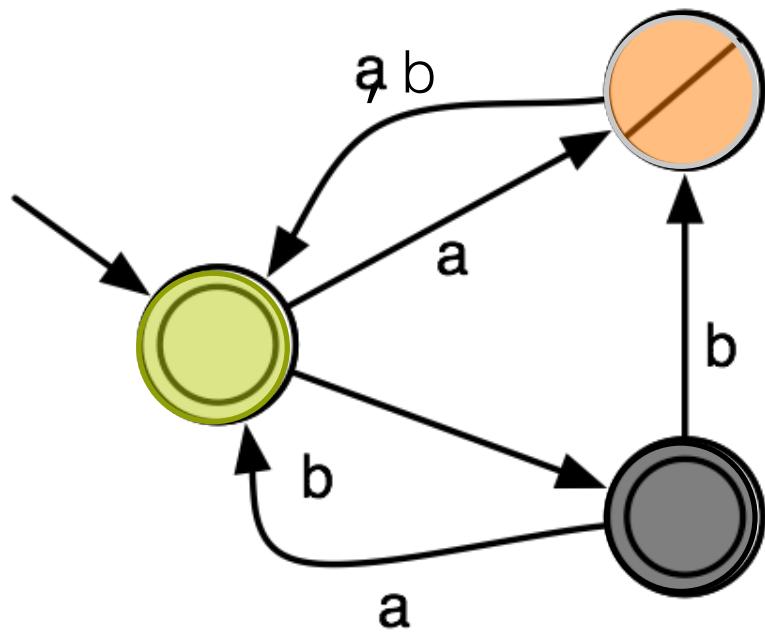
	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	1	0	1	1	1	1	0	0	1	0
a	0	1	1	0	1	0	1	1	1	1
b	1	1	0	0	1	1	1	1	1	1
aa	1	0	1	1	1	1	0	0	1	0
ab	1	0	1	1	1	1	0	0	1	0
ba	1	0	1	1	1	1	0	0	1	0
bb	0	1	1	0	1	0	1	0	1	1
aaa	0	1	1	0	1	0	1	1	1	1
aab	1	1	0	0	1	1	1	1	1	1
aba	0	1	1	0	1	0	1	1	1	1

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	1	0	1	1	1	1	0	0	1	0
a	0	1	1	0	1	0	1	1	1	1
b	1	1	0	0	1	1	1	1	1	1
aa	1	0	1	1	1	1	0	0	1	0
ab	1	0	1	1	1	1	0	0	1	0
ba	1	0	1	1	1	1	0	0	1	0
bb	0	1	1	0	1	0	1	1	1	1
aaa	0	1	1	0	1	0	1	1	1	1
aab	1	1	0	0	1	1	1	1	1	1
aba	0	1	1	0	1	0	1	1	1	1

Answer



Answer



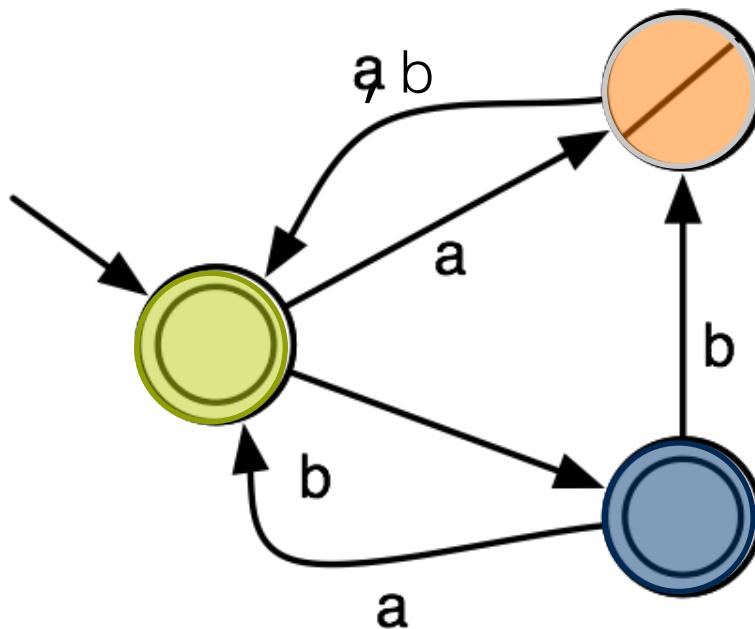
	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	1	0	1	1	1	1	0	0	1	0
a	0	1	1	0	1	0	1	1	1	1
b	1	1	0	0	1	1	1	1	1	1
aa	1	0	1	1	1	1	0	0	1	0
ab	1	0	1	1	1	1	0	0	1	0
ba	1	0	1	1	1	1	0	0	1	0
bb	0	1	1	0	1	0	1	0	1	1
aaa	0	1	1	0	1	0	1	1	1	1
aab	1	1	0	0	1	1	1	1	1	1
aba	0	1	1	0	1	0	1	1	1	1

Look at colors of Hankel Matrix rows

—
a
b
aa
ab
ba
bb
aaa
aab
aba

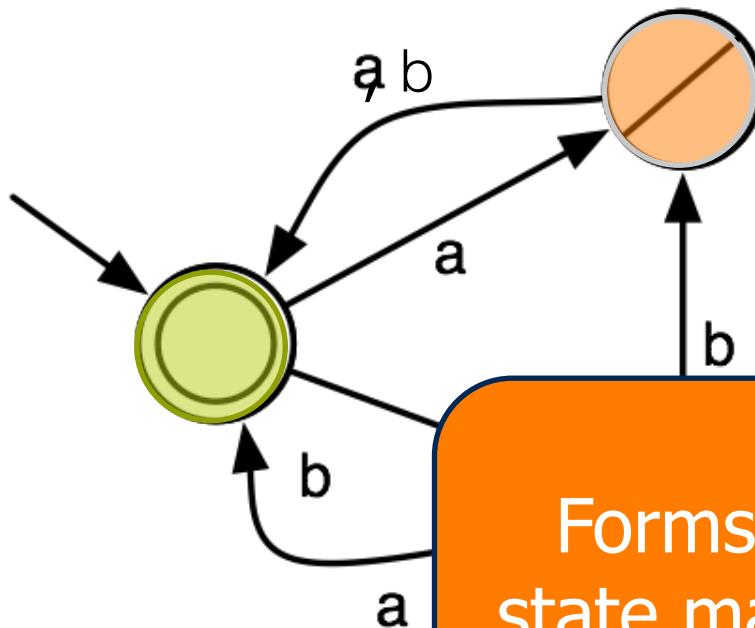
–
a
b
aa
ab
ba
bb

access strings and *one-letter extensions*
provide sufficient information



—
a
b
aa
ab
ba
bb

access strings and *one-letter extensions*
provide sufficient information



Forms the basis for
state machine learning
algorithms

STATE MACHINE LEARNING

(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

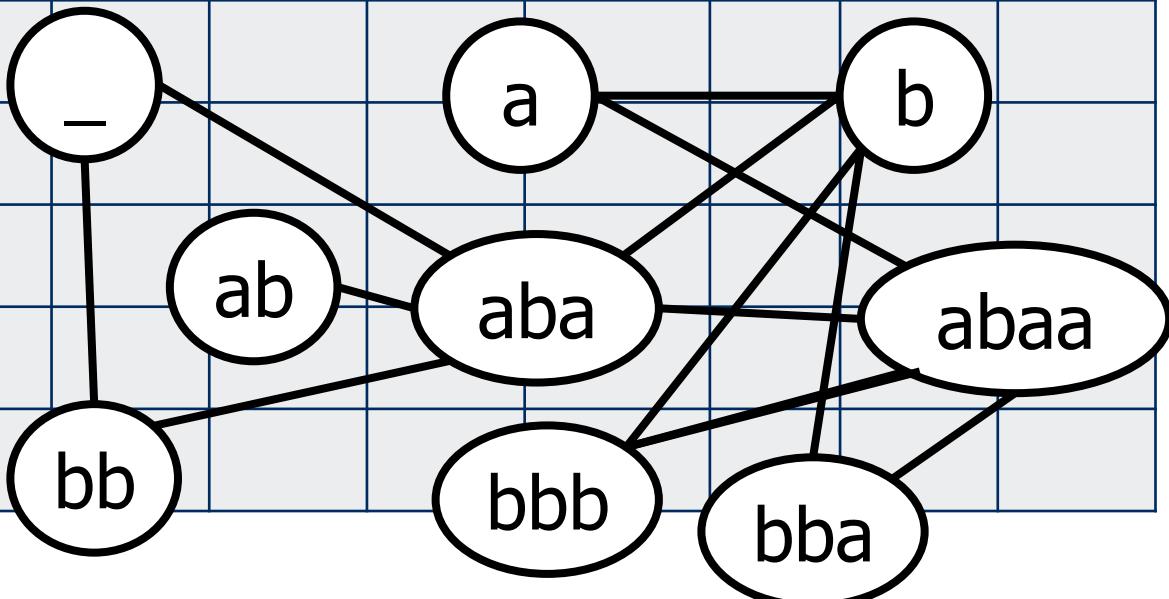
	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

$(a,+)$ $(b,-)$ $(aba,+)$ $(abaa,-)$ $(bbb,+)$ $(bba,+)$

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										



(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

a b ba aba aa

baa abaa

bb bbb bba

0 1 1

0 1

ab 1

0

aba 1 0

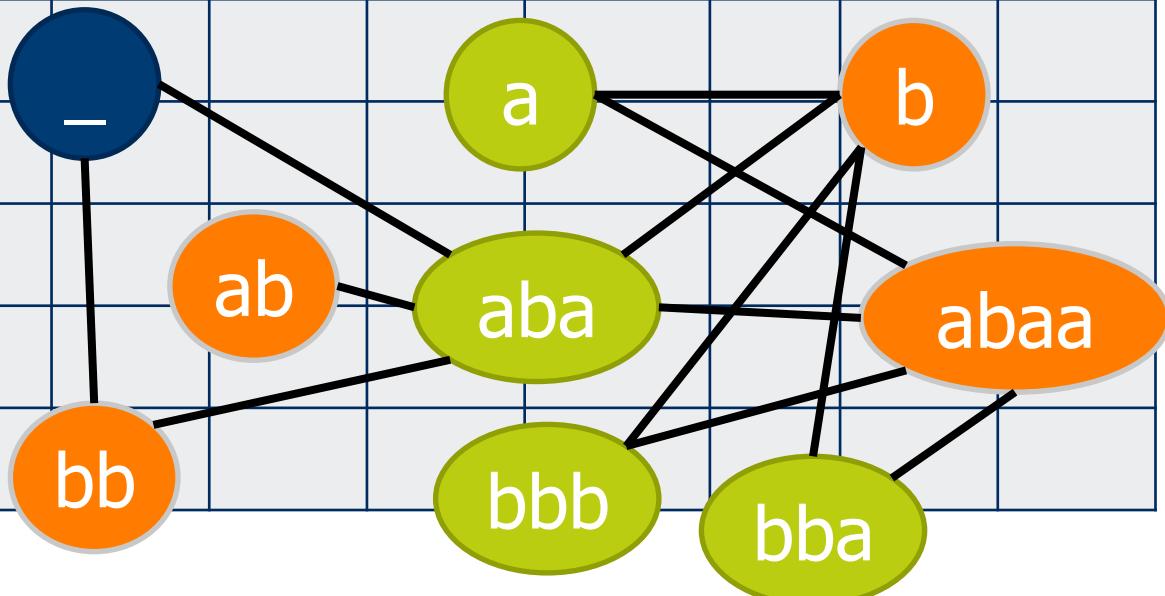
abaa 0

bb 1 1

bbb 1

bba 1

Try to make state machine...



(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

a b ba aba aa

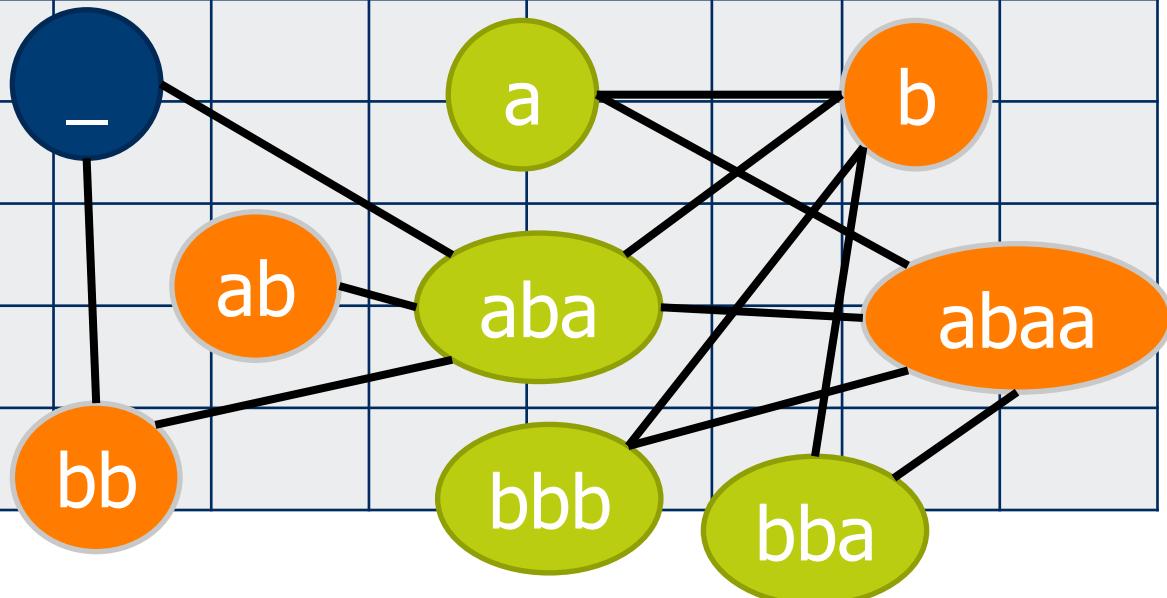
baa abaa bb bbb bba

Try to make state machine...

It is **inconsistent!**

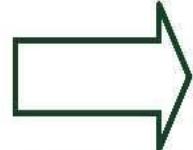
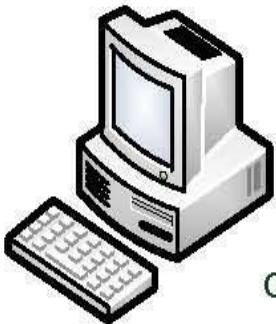
(b and bb end in same state,
but bbb does not)

ab	a	b	ba	aba	aa
aba	1	0			
abaa	0				
bb		1	1		
bbb	1				
bba	1				



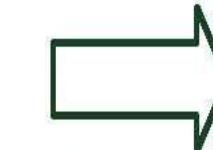
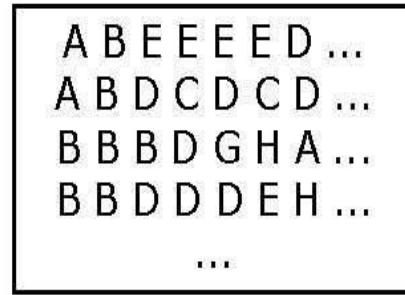
We need more information...

software system

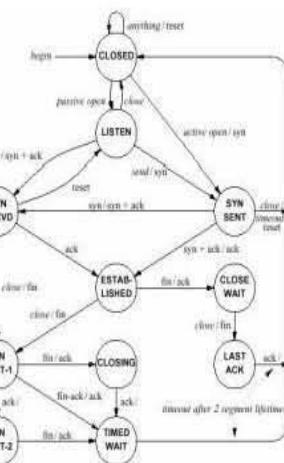


system call or
communication logs

execution traces



state machine
learning



software
model

Queries (input)



Active learning

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

What to ask?

Ask: bb

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

Ask: bb

	_	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
_	1	0		1				0		1	1
a	1			1			0				
b	0		1	1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb	1	1	1								
bbb	1										
bba	1										

If + _ distinguishes between b and bb

Ask: bb

	-	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-	1	0		1				0		1	1
a	1			1			0				
b	0		0	1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb	0	1	1								
bbb	1										
bba	1										

If - b distinguishes between b and bb

Active learning with L*

- Aims to discover:
 - **access strings**, smallest strings for every color
 - **one-letter extensions**, and
 - their **distinguishing suffixes**
- Maintains a distinguishing table
- Which distinguishing suffixes is determined during learning

	—	a
—	1	0
a	0	1
b	1	1
aa	1	0
ab	1	0
ba	1	0
bb	0	1

Active learning with L*

- Start with an empty table, with one access string and one distinguishing extension:
 - _ and _
- Ask membership queries for _ and its one letter extensions:
 - _ in L? \rightarrow yes
 - a in L? \rightarrow no
 - b in L? \rightarrow yes

	-
-	1
a	0
b	1

Active learning with L^*

- Can we construct a state machine?
 - No!
 - a points to a new state
- Ask membership queries for its one-letter extensions:
 - aa in L ? \rightarrow yes
 - ab in L ? \rightarrow yes

—	—
—	1
a	0
b	1
aa	1
ab	1

Active learning with L^*

- Can we construct a state machine?
 - Yes!
- Ask an equivalence query:
 - $L(M) = L? \rightarrow$ no, M for instance accepts `aabbaa`

—	—
—	1
a	0
b	1
aa	1
ab	1

Active learning with L^*

- $C = aabbaa, -$
- Where does M make a mistake?
 - *Replace C prefixes with reached access strings*
- Ask membership queries:
 - $_bbaa$ in L ? \rightarrow no
 - $_baa$ in L ? \rightarrow no
 - $_aa$ in L ? \rightarrow yes

	-
-	1
a	0
b	1
aa	1
ab	1

Active learning with L*

- $C = \text{aabbaa}, -$
- Where does M make a mistake?
 - *Replace C prefixes with reached access strings*
- Ask membership queries:
 - $_ \text{bbaa}$ in L ? \rightarrow no
 - $_ \text{baa}$ in L ? \rightarrow no
 - $_ \text{aa}$ in L ? \rightarrow

	-
-	1
a	0
b	1
aa	1
ab	1

aa is a new
distinguishing extension!

Active learning with L*

- Add the column aa
- Ask membership queries to complete the table

	–	aa
–	1	1
a	0	0
b	1	0
aa	1	1
ab	1	1

Active learning with L*

- Add the column aa
- Ask membership queries to complete the table
- b points to a new state
- Add its one-letter extensions, ask membership queries to complete

	—	aa
—	1	1
a	0	0
b	1	0
aa	1	1
ab	1	1
ba	1	1
bb	0	0

Active learning with L*

- Add the column aa
- Ask membership queries

The table defines a state machine, ask an equivalence query → yes!

Done!

complete

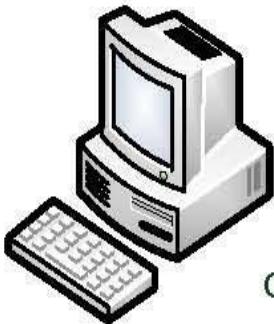
	—	aa
—	1	1
a	0	0
b	1	0
aa	1	1
ab	1	1
ba	1	1
bb	0	0

Active learning with L*

1. Initialize the distinguishing table T
2. Complete T until it defines a state machine M
3. While(Eq(M) gives a counterexample c)
 1. Use Mem() and c to find a new distinguishing extension d
 2. Add d to T
 3. Complete T until it defines a state machine M

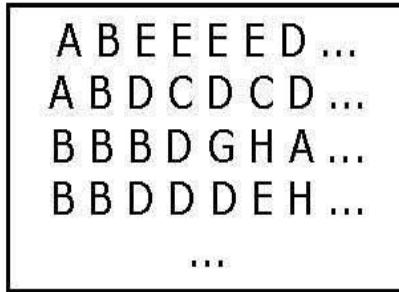
Use the information we have...

software system

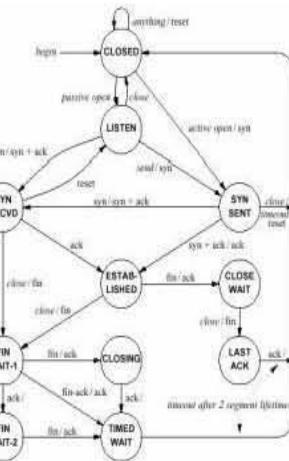


system call or
communication logs

execution traces



state machine
learning



software
model

Passive learning

Key issue: coloring leads to inconsistencies

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

```

graph LR
    _(( )) --> a((a))
    a --> b((b))
    b --> ab((ab))
    ab --> aba((aba))
    aba --> aa((aa))
    aa --> baa((baa))
    baa --> abaa((abaa))
    bb((bb)) --> bba((bba))
    bba --> bbb((bbb))
    bbb --> b((b))
    b --> ab((ab))
    ab --> aba((aba))
    aba --> baa((baa))
    baa --> abaa((abaa))
    abaa --> bba((bba))
    bba --> bbb((bbb))
    baa -.-> abaa((abaa))
    abaa -.-> bba((bba))
    bba -.-> bbb((bbb))
    
```

Solution: state merging

- Very effective **greedy** method
 - Initially assume a large DFA
 - Iteratively combine similar states
 - Until no more state can be combined
- State-of-the-art since 1998

color iteratively, copy values, force determinism!

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

color iteratively, copy values, force determinism!

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

color iteratively, copy values, force determinism!

	_	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

can a be merged with _?

color iteratively, copy values, force determinism!

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

if a and _ are the same,
ab and b are the same

color iteratively, copy values, force determinism!

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

a and _ can be merged,
but there is 0 overlap in values

evidence suggests not to merge

color iteratively, copy values, force determinism!

	_	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

can b be merged with _?

color iteratively, copy values, force determinism!

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

no!

color iteratively, copy values, force determinism!

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

color iteratively, copy values, force determinism!

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

there is evidence suggesting
ab and _ to be the same

color iteratively, copy values, force determinism!

	-	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

even more in determinization

color iteratively, copy values, force determinism!

	-	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-		1	0		1	0		0		1	1
a	1	0		1			0				
b	0			1					1		
ab		1	0		1	0		0		1	1
aba	1	0		1			0				
abaa	0										
bb		1	1								
bbb	1										
bba	1										

merge and copy values!

color iteratively, copy values, force determinism!

	-	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-		1	0		1	0		0		1	1
a	1	0		1			0				
b	0			1					1		
ab		1	0		1	0		0		1	1
aba	1	0		1			0				
abaa	0			1					1		
bb		1	1								
bbb	1										
bba	1										

continue merging
maximizing evidence

color iteratively, copy values, force determinism!

	-	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-		1	0		1	0		0		1	1
a	1	0		1			0				
b	0			1					1		
ab		1	0		1	0		0		1	1
aba	1	0		1			0				
abaa	0			1					1		
bb		1	1								
bbb	1	0		1			0				
bba	1	0		1							

continue merging
maximizing evidence

Red-blue fringe state merging

- *To improve speed, use trees!*
1. Start from a **prefix tree acceptor**
 2. Color the initial state red, and its children blue
 3. For all **consistent** red-blue state pairs R-B
 1. **Merge** B with R
 2. Compute evidence **score**
 3. **Undo merge**
 4. If there is a high scoring and consistent R-B merge
 1. **Merge** B with R
 5. Else
 1. Color B with most occurrences (or most shallow first, ...) red
 6. Color all children of red states blue
 7. Goto 3

Red-blue fringe state merging

- *To improve speed, use trees!*
 1. Start from a **prefix tree acceptor**
 2. Color the initial state red, and its children blue
 3. For all **consistent** red-blue state pairs R-B
 1. **Merge** B with R
 2. Compute evidence **score**
 3. **Undo merge**
 4. If there is a high scoring and consistent R-B merge
 1. **Merge** B with R
 5. Else
 1. Color B with most occurrences (or most shallow first, ..) red
 6. Color all children of red states blue
 7. Goto 3

efficient active learning
algorithms use trees as well

or... solve the NP-hard problem by SAT

- Generate the prefix-tree T
- Initialize the set of colors L
- Construct a SAT formula F using T and L
- Solve F using a SAT solver
- If F is unsatisfiable add a color to L , goto 3
- Return the DFA found in step 4

Satisfiability encoding

Variables	Range	Meaning
$x_{v,i}$	$v \in V; i \in C$	$x_{v,i} \equiv 1$ iff state v has color i
$y_{a,i,j}$	$a \in \Sigma; i, j \in C$	$y_{a,i,j} \equiv 1$ iff parents of states with color j and incoming label a have color i
z_i	$i \in C$	$z_i \equiv 1$ iff an accepting state has color i
Clauses	Range	Meaning
$(x_{v,1} \vee x_{v,2} \vee \dots \vee x_{v,k})$	$v \in V$	every state has at least one color
$(\neg x_{v,i} \vee z_i) \wedge (\neg x_{w,i} \vee \neg z_i)$	$v \in V_+; w \in V_-; i \in C$	accepting states cannot have the same color as rejecting states
$(y_{l(v),i,j} \vee \neg x_{p(v),i} \vee \neg x_{v,j})$	$v \in V; i, j \in C$	a parent relation is set when a state and its parent are colored
$(\neg y_{a,i,h} \vee \neg y_{a,i,j})$	$a \in \Sigma; h, i, j \in C; h < j$	each parent relation can target at most one color
Redundant Clauses	Range	Meaning
$(\neg x_{v,i} \vee \neg x_{v,j})$	$v \in V; i, j \in C; i < j$	every state has at most one color
$(y_{a,i,1} \vee y_{a,i,2} \vee \dots \vee y_{a,i,k})$	$a \in \Sigma; i \in C$	each parent relation must target at least one color
$(\neg y_{l(v),i,j} \vee \neg x_{p(v),i} \vee x_{v,j})$	$v \in V; i, j \in C$	a parent relation forces a state once the parent is colored
$(\neg x_{v,i} \vee \neg x_{w,i})$	$i \in C; (v, w) \in E$	all determinization conflicts explicitly added as clauses

Replacing Myhill-Nerode by Markov

- Two rows r, r' can have the same color if
 - $r = c * r'$, c a constant
 - i.e., *if the suffix probability is independent of the prefix!*
- Given an **unlabeled dataset**
 - Put the sample probabilities in a Hankel matrix
 - Determine the minimal coloring of the matrix
- Using, i.e.:
 - State merging (ALERGIA)
 - Singular value decomposition (**spectral learning**)
 - satisfiability or graph coloring
 - ...

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	1	0	1	1	1	1	0	0	1	0
a	0	1	1	0	1	0	1	1	1	1
b	1	1	0	0	1	1	1	1	1	1
aa	1	0	1	1	1	1	0	0	1	0
ab	1	0	1	1	1	1	0	0	1	0
ba	1	0	1	1	1	1	0	0	1	0
bb	0	1	1	0	1	0	1	0	1	1
aaa	0	1	1	0	1	0	1	1	1	1
aab	1	1	0	0	1	1	1	1	1	1
aba	0	1	1	0	1	0	1	1	1	1

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0.46	0	0.08	0.08	0.03	0.07	0	0	0.01	0
a	0	0.08	0.03	0	0.01	0	0	0.01	0	0.01
b	0.08	0.07	0	0	0.01	0.03	0.01	0.01	0	0.01
aa	0.08	0	0.01	0.01	0	0.01	0	0	0	0
ab	0.03	0	0	0	0	0	0	0	0	0
ba	0.07	0	0.01	0.01	0	0.01	0	0	0	0
bb	0	0.03	0.01	0	0	0	0	0	0	0
aaa	0	0.01	0	0	0	0	0	0	0	0
aab	0.01	0.01	0	0	0	0	0	0	0	0
aba	0	0	0	0	0	0	0	0	0	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0.46	0	0.08	0.08	0.03	0.07	0	0	0.01	0
a	0	0.08	0.03	0	0.01	0	0	0.01	0	0.01
b	0.08	0.07	0	0	0.01	0.03	0.01	0.01	0	0.01
aa	0.08	0	0.01	0.01	0	0.01	0	0	0	0
ab	0.03	0	0	0	0	0	0	0	0	0
ba	0.07	0	0.01	0.01	0	0.01	0	0	0	0
bb	0	0.03	0.01	0	0	0	0	0	0	0
aaa	0	0.01	0	0						
aab	0.01	0.01	0	0						
aba	0	0	0	0						

Copying values is less important

Merged rows are statistically close anyway...

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0.46	0	0.08	0.08	0.03	0.07	0	0	0.01	0
a	0	0.08	0.03	0	0.01	0	0	0.01	0	0.01
b	0.08	0.07	0	0	0.01	0.03	0.01	0.01	0	0.01
aa	0.08	0	0.01	0.01	0	0.01	0	0	0	0
ab	0.03	0	0	0	0	0	0	0	0	0
ba	0.07	0	0.01	0.01	0	0.01	0	0	0	0
bb	0	0.03	0.01	0	0	0	0	0	0	0
aaa	0	0.01	0	0						
aab	0.01	0.01	0	0						
aba	0	0	0	0						

The Hankel matrix is fully specified given the input data!

This is used by spectral learning

Spectral learning

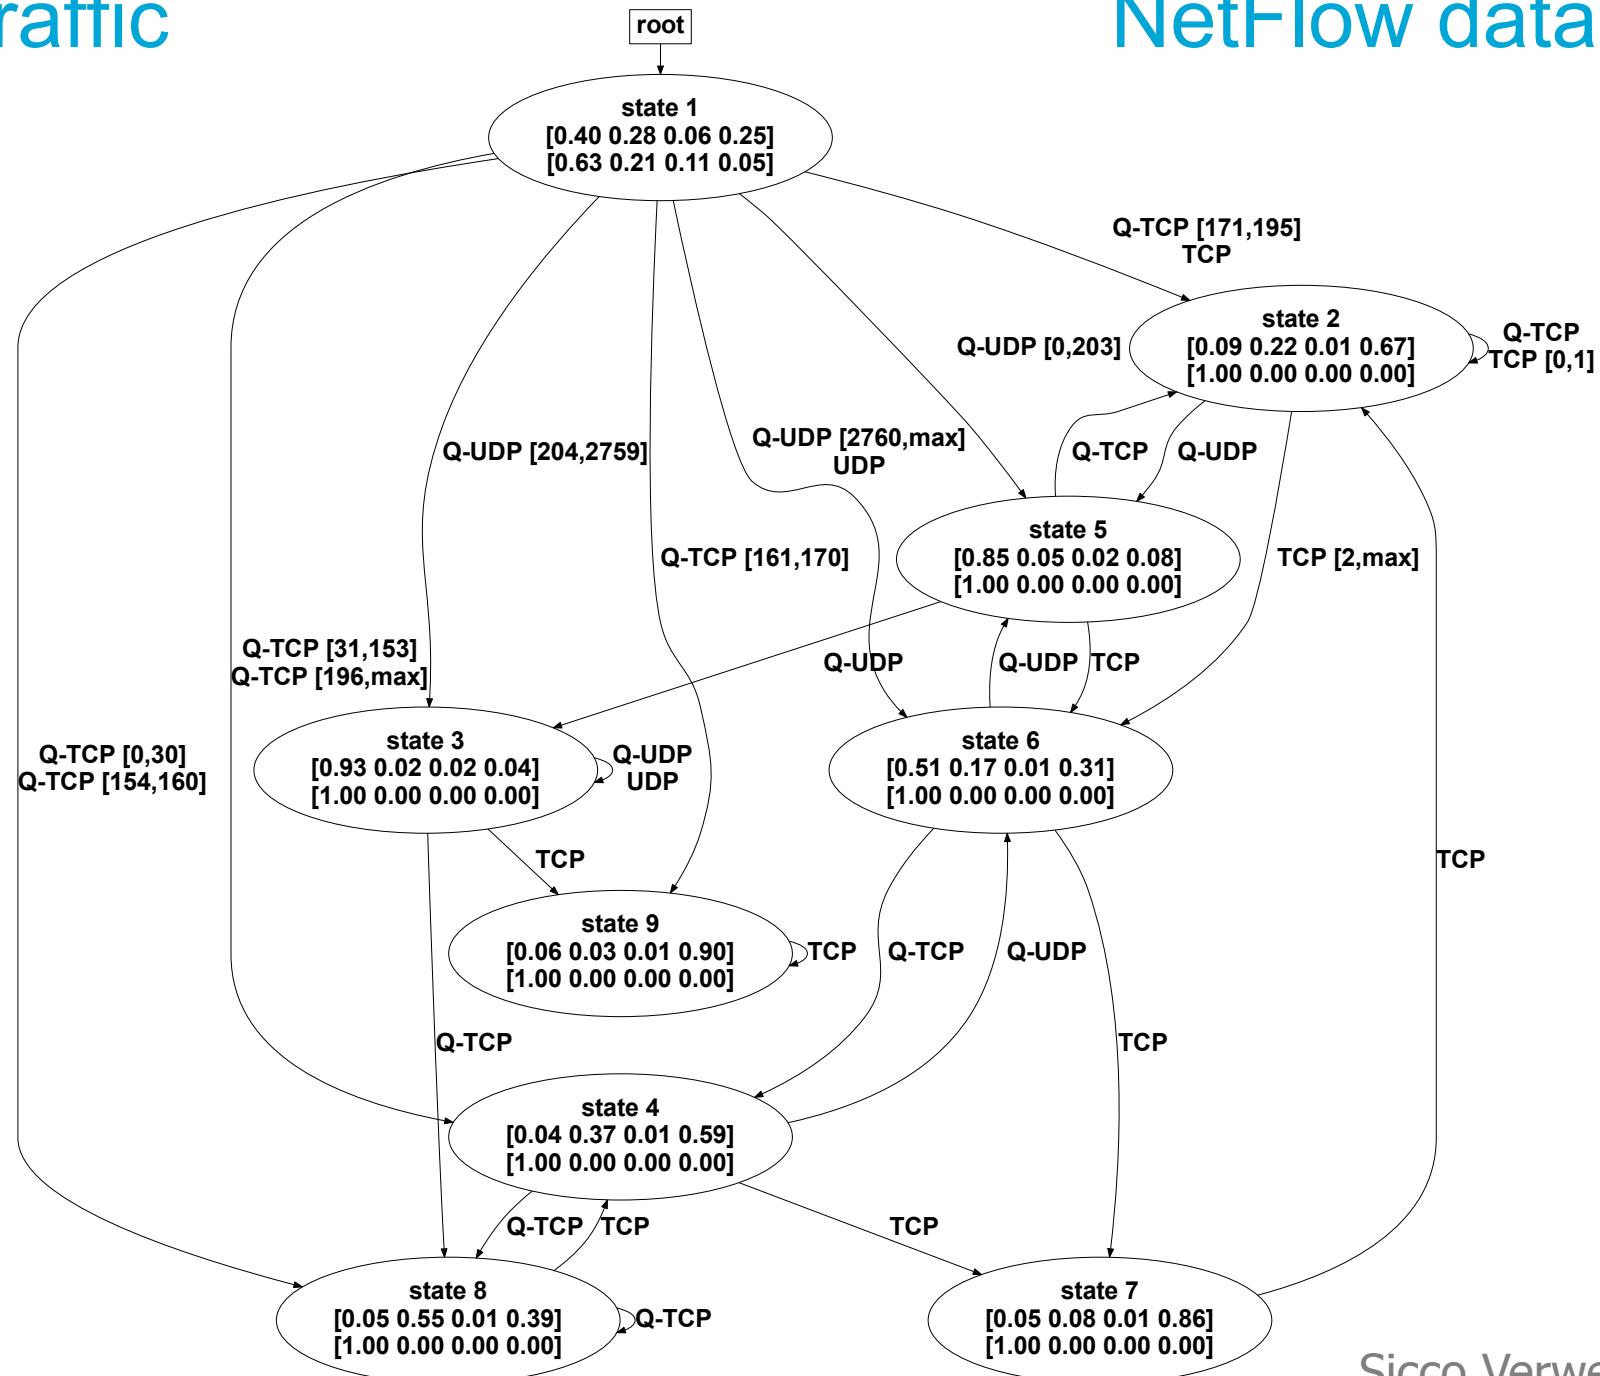
- Problem: the matrix M is fully specified, but very noisy
- Three steps:
 1. Compute singular values of M
 2. Keep the n largest singular values, recompute M
 3. Color M assuming it is exact
- Step 2 compresses M to an n -state object, removing noise
 - Given sufficient data, this turns out to be a non-deterministic PFA!
 - If not, what is learned it is not really defined...
- Step 3 is not hard because M is fully specified

Key ingredients

- When are two states **inconsistent**? (both active and passive)
 - Positive gets merged with negative
 - Input/output does not match
 - Other events are possible
 - Statistics do not match
 - If the Markov property is violated
 - ...
- What to use as merge **score**? (passive)
 - Number of positive-positive, negative-negative merges
 - Overlap in sets of possible events
 - Statistical distance
 - ...
- Should we ensure **correctness** at all? (spectral, graph coloring)

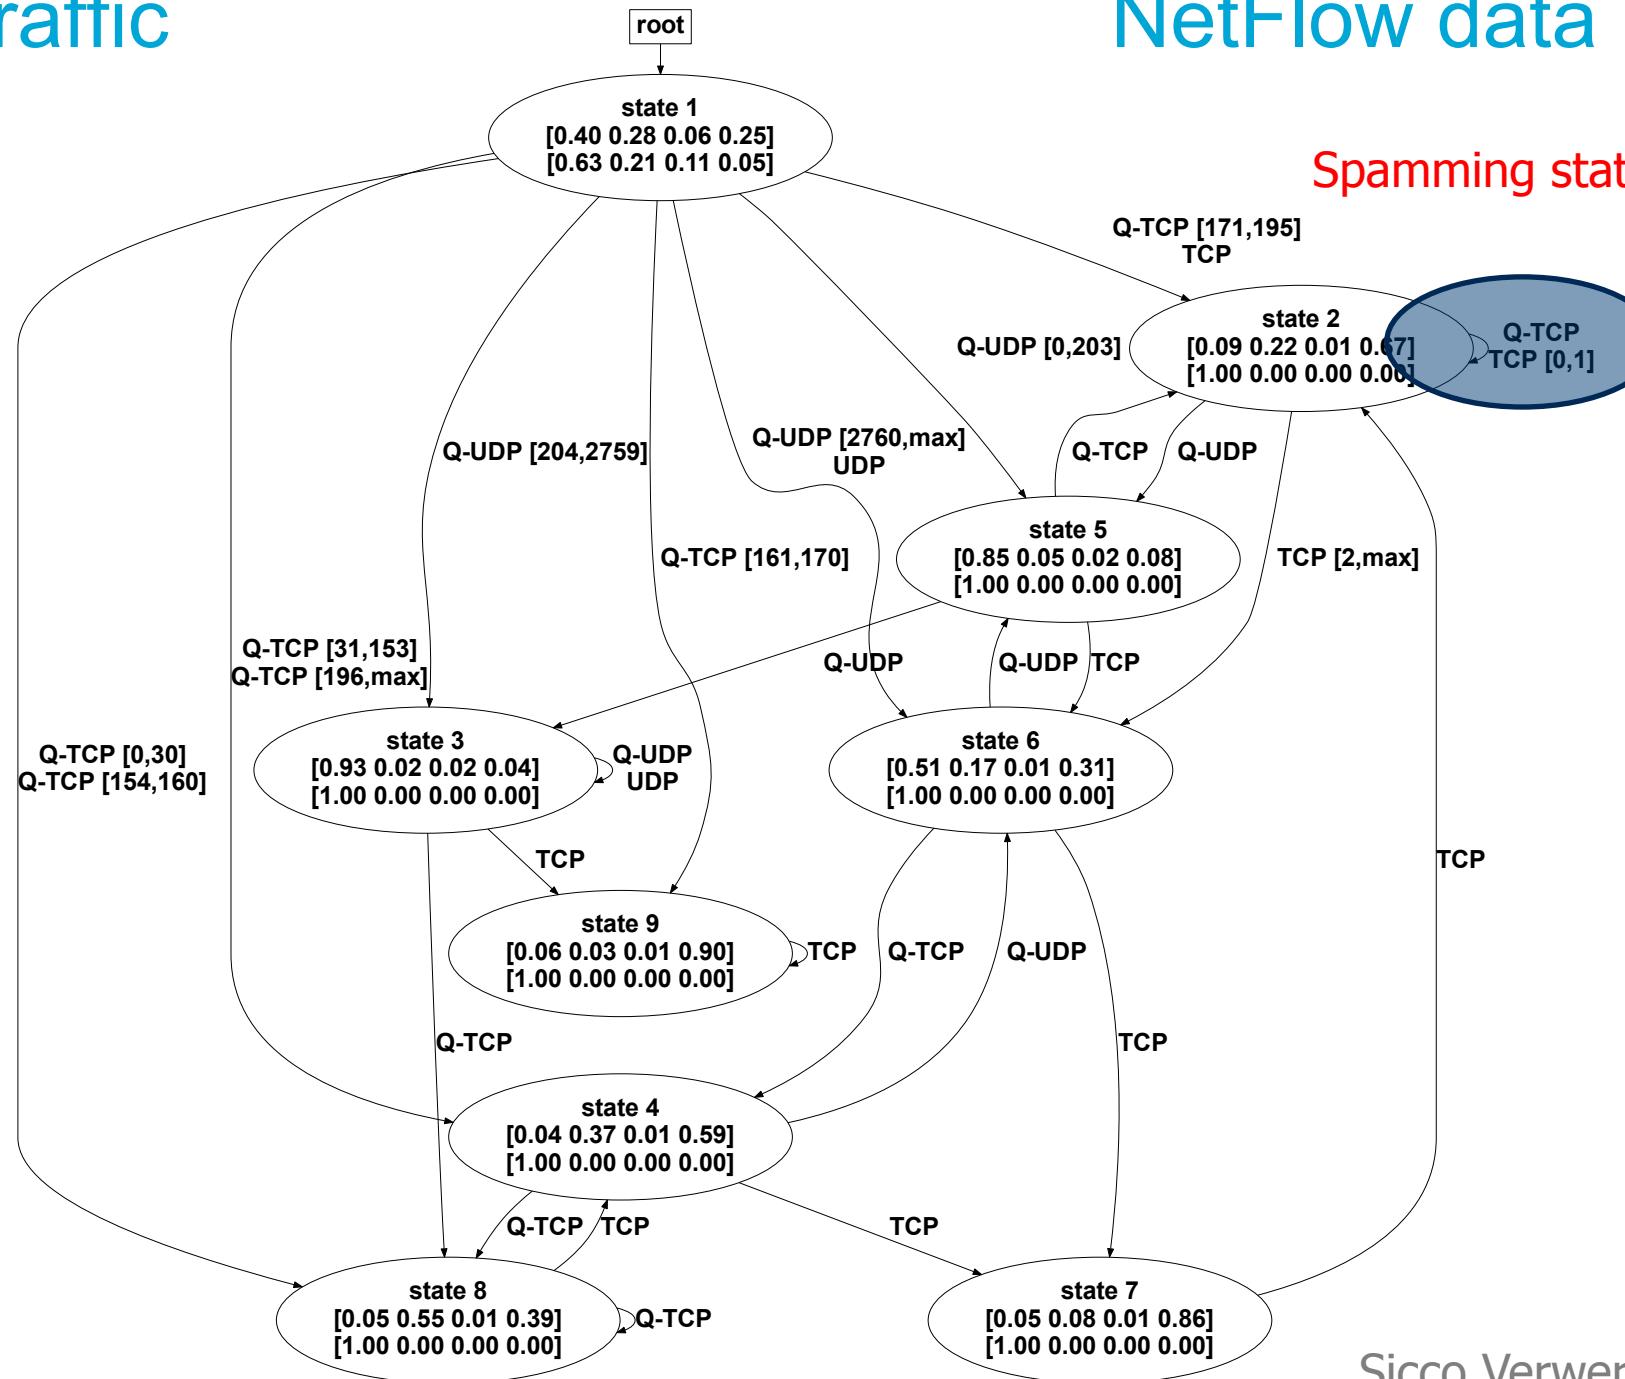
Botnet traffic

NetFlow data

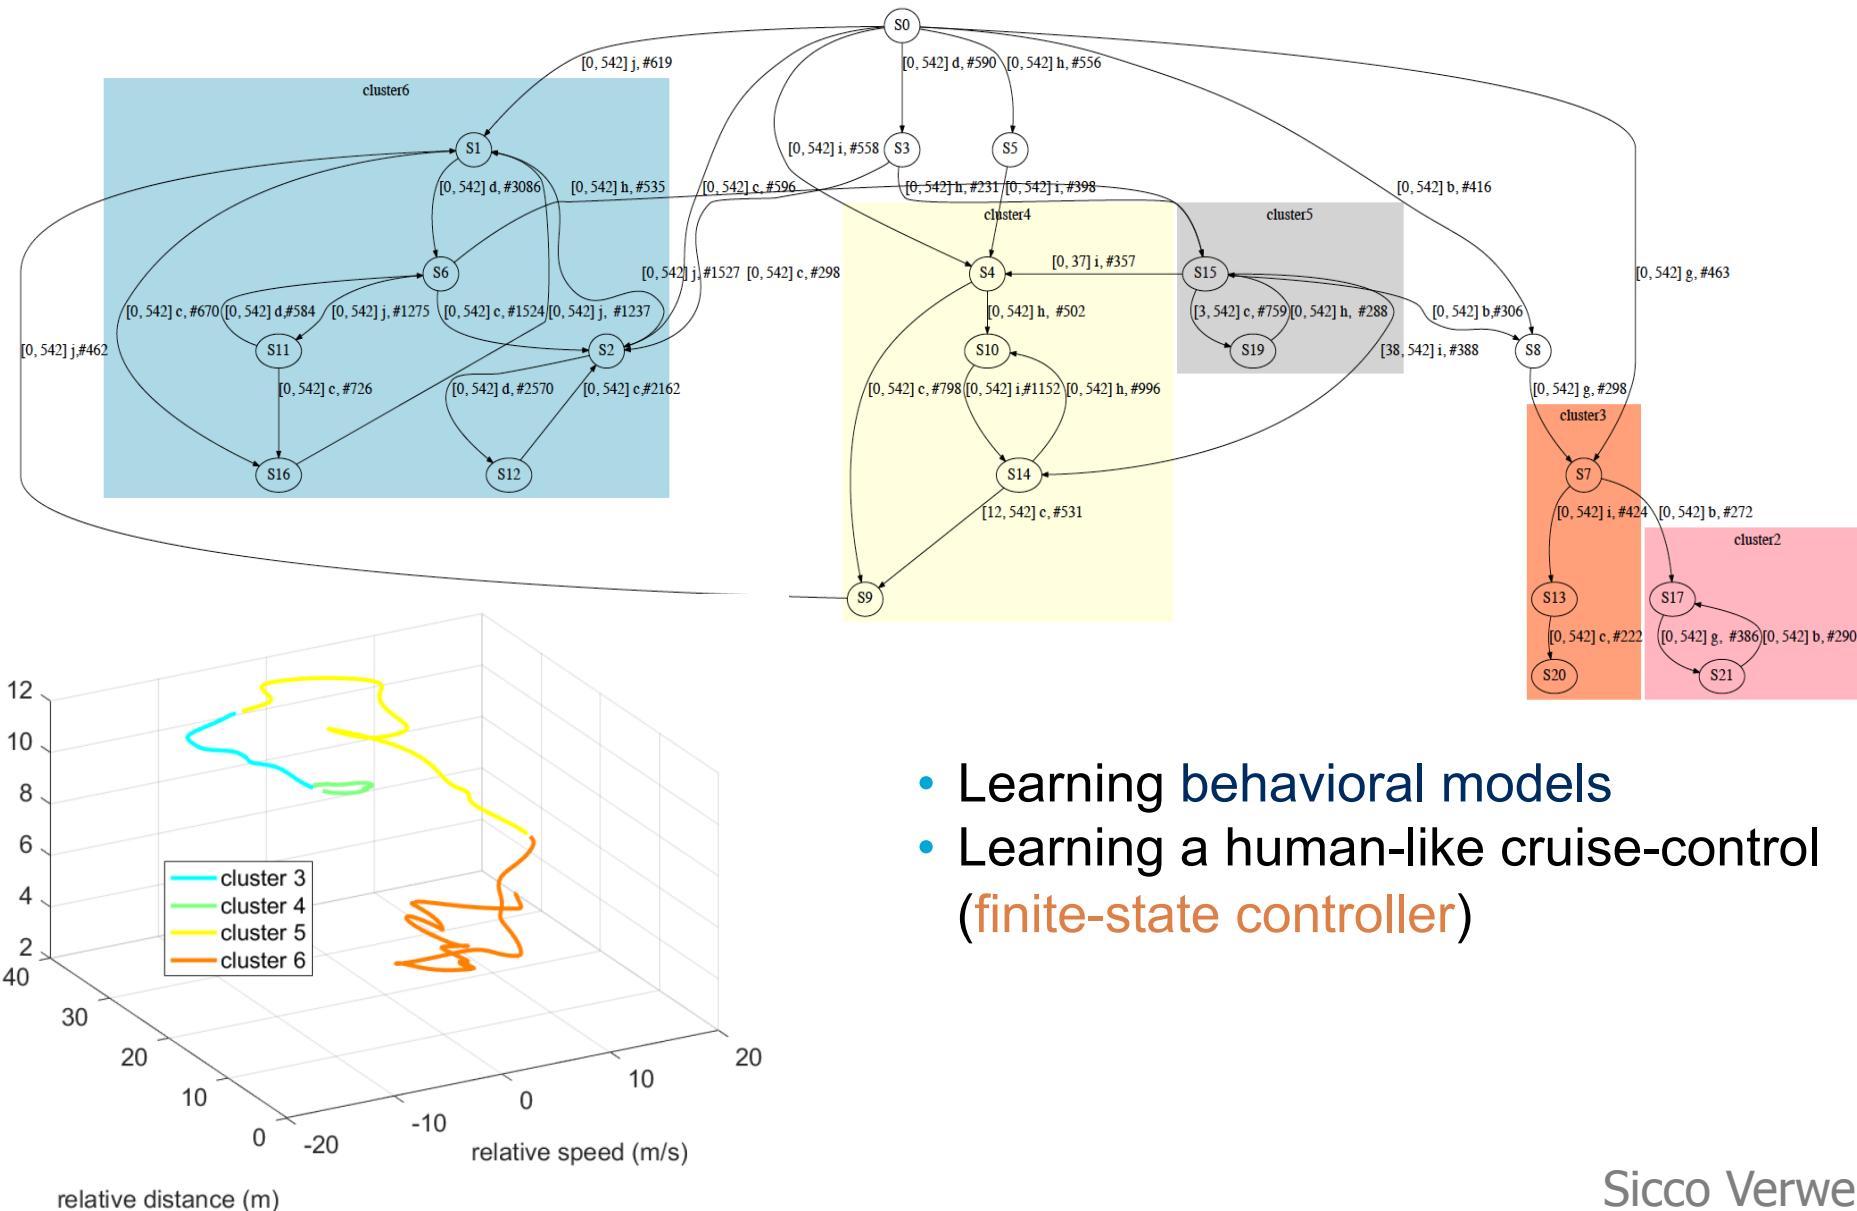


Botnet traffic

NetFlow data



Automated driving



- Learning behavioral models
- Learning a human-like cruise-control (finite-state controller)

Near Future

- First method for systematically analyzing network data
 1. Connect to a monitor
 2. Leave alone for a week
 3. Analyze the learned state machines
- Develop learning tool

fast & flexible

